

Content

■ 3.1 Base Software Stack

- Foundational Software Stack
- Environment Modules System
- Compiling and Running C Programs

■ 3.2 Workload Management with SLURM

- The Essentials of SLURM
- Job Directives
- Computing Resource Allocation

■ 3.3 Getting Started with Docker Containers

- Docker basics
- Launching a Jupyter Notebook server inside a Docker

■ 3.4 Containerization with Singularity

- Containers in Supercomputing
- Building Containers



Getting Started with Docker Container

Dockers Basics

- **Docker = leading container platform**
 - `https://www.docker.com`
- **Packages app + dependencies → standardized unit**
- **Runs the same across Linux, Windows, macOS**
- **Containers = isolation from environment differences**
- **Benefits: reproducibility, portability, ...**

Docker Image

- **Lightweight, standalone, executable package**
- **Includes:**
 - **Code**
 - **Runtime**
 - **Libraries**
 - **System tools & settings**
- **Ensures identical behavior across environments**

Docker Hub

- **Official repository for container images**
 - `https://hub.docker.com`
- **Public & private images supported**
- **Developers can share or reuse images**
- **Example: `jorditorresbcn/dl`**

Push Image to Docker Hub

1. Create account on Docker Hub

2. Log in:

```
docker login
```

3. Build image locally:

```
docker build -t <user>/<image>:<tag> .
```

4. Push to Docker Hub:

```
docker push <user>/<image>:<tag>
```

5. Others can pull it:

```
docker pull <user>/<image>:<tag>
```

Install Docker

- **Windows:**

- Docker Desktop (includes CLI, Compose, Kubernetes)

- **Mac:**

- Docker Desktop (\geq macOS 10.14)

- **Linux:**

- Install Docker Engine via package manager (e.g., Ubuntu, CentOS, Fedora)

- **Verify installation:**

```
docker --version
```

```
docker run hello-world
```

Download Docker Image

- **Command:**

```
docker pull jorditorresbcn/dl
```

- **Verify:**

```
docker images
```

- **Info: name, tag, size**

- **Benefit: reuse pre-built images instead of building from scratch**

Run & Stop Containers

- **Run basic:**

```
docker run <image>
```

- **Run interactive:**

```
docker run -it jorditorresbcn/dl
```

- **List running:**

```
docker ps
```

- **Stop / kill:**

```
docker stop <id_or_name>
```

```
docker kill <id_or_name>
```

- **Restart / remove:**

```
docker start <id_or_name>
```

```
docker rm <id_or_name>
```

Tasks

- **Task 3.4 – Install Docker**
 - Install on your platform, verify with `docker --version`
 - Test with `docker run hello-world`
- **Task 3.5 – Download Image**
 - Pull `jorditorresbcn/dl`
 - Verify with `docker images`
- **Task 3.6 – Run Image**
 - Run interactive container: `docker run -it jorditorresbcn/dl`
 - Inside: check OS (`cat /etc/os-release`), list Python libs (`pip list`)
- **Task 3.7 – Stop Container**
 - Exit with `exit` or stop with `docker stop <id>`




Launching a Jupyter Notebook Server Inside a Docker Container

Port Mapping in Docker

```
docker run -it -p 8888:8888 --name test  
jorditorresbcn/dl:latest
```

- **-p 8888:8888** → maps container port → host port
- **Access service inside container (e.g., Jupyter)**
from browser: `http://localhost:8888`
- **Ports = endpoints for services, allow multiple apps without conflict**



Jupyter Notebook Basics


- **Open-source web application for creating and sharing documents**
- **Combines in a single file:**
 - Live executable code (Python), Narrative text (Markdown), Equations (LaTeX), Visualizations and interactive widgets, ..
- **Notebook structure:**
 - Organized in cells → run independently
 - Immediate feedback → great for teaching & experimentation
- **File format:**
 - `.ipynb` (JSON with code, text, outputs, metadata)
- **Where we use them in this course:**
 - Docker
 - Google Colab
 - MareNostrum 5

Starting Jupyter in the Container

■ Command:

```
jupyter notebook --ip=0.0.0.0 --allow-root
```

- `--ip=0.0.0.0` → accept connections from any IP (not just localhost)
- Defaults to port 8888 (or next available if in use)
- `--allow-root` → required in container environments



Accessing Jupyter Notebook

- **Start server with:**

```
jupyter notebook --ip=0.0.0.0 --port=8888 --no-browser
```

- **Terminal shows URL with token → copy/paste in host browser**

- **Default access:** `http://127.0.0.1:8888`
(password: dl)

- **Now you can create and run notebooks from your browser**

Tasks

- **Task 3.8 – Run Docker with Port Mapping → run container with -p 8888:8888 and test access**
- **Task 3.9 – Start Jupyter Server → run jupyter notebook inside container**
- **Task 3.10 – Create Test Notebook → confirm setup with `print("Docker and Jupyter are working!")`**



Containers in Supercomputing

Containers in Supercomputing

- In HPC, reproducibility and portability are critical.
- Containers package the complete software stack (OS, libraries, dependencies, tools).
- This ensures consistent execution across different clusters, avoiding local configuration issues.
- Increasingly important for complex scientific workflows and ML applications.

Docker vs Singularity in HPC

■ Docker:

- Industry standard for cloud & microservices.
- Requires root privileges → **unsafe in shared systems**.
- Not designed for MPI tightly coupled computations.

■ Singularity:

- Created by the HPC community.
- **Runs without root access**, integrates with SLURM.
- Provides direct access to GPUs, filesystems, and interconnects.
- De facto standard for HPC at BSC (SingularityCE).

Running Singularity on MareNostrum 5

- **Load module:** `module load SINGULARITY/3.11.5`
- **Container format:** `.sif` **(Singularity Image Format)**
- **Basic execution with GPU passthrough:**
`singularity exec --nv image.sif python train.py`
- **BSC helper tool:** `bsc_singularity`
 - **Automatically adds GPU support and filesystem bindings**
 - **Provides easy commands:**
`bsc_singularity ls`
`bsc_singularity exec <container>`
`bsc_singularity shell <container>`



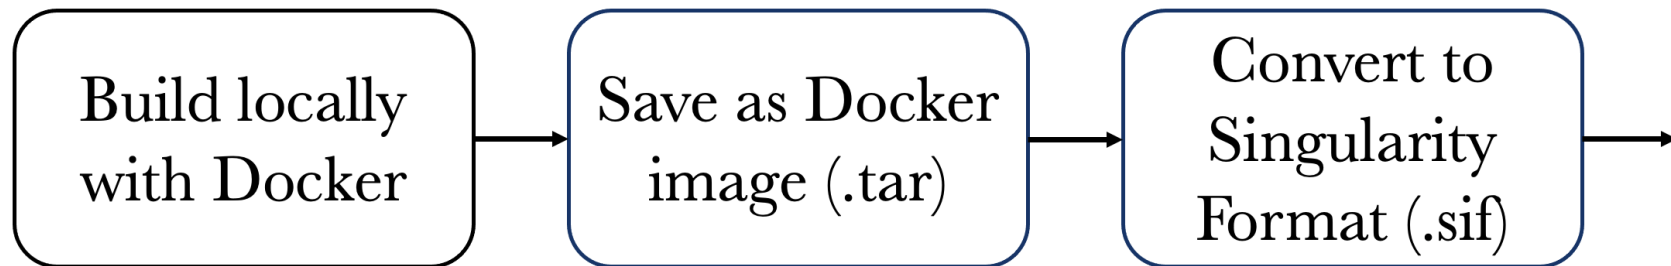
Building Containers

Overview

- **Containers can be built from different sources:**
 - Singularity Library
 - Docker Hub
 - Definition file
 - NVIDIA NGC Docker images
- **The common workflow:**
 - Build locally (requires root + internet).
 - Transfer to HPC system (e.g., MareNostrum 5).

Preferred hybrid approach

- develop with Docker, deploy with Singularity:



- Ensures flexibility in development + compatibility in HPC.
- For more detail visit section “Building Containers”.

Pb: Containers

■ **Tasks included:**

Task 3.4 – Install Docker in Your Platform

Task 3.5 – Download Docker Image

Task 3.6 – Run Docker Image

Task 3.7 – Stop a Docker Container

Task 3.8 – Run Docker with Port Mapping

Task 3.9 – Start the Jupyter Notebook Server

Task 3.10 – Create and Run a Test Notebook

■ **Deliverable:**

- Upload a single PDF (per group) to the intranet racó@FIB containing one slide per task. Each slide should report results (if applicable) or briefly explain how the task was completed.
- In class (evaluation day), one student (chosen at random) will give a short “*elevator pitch*”-style presentation — clear, concise, and straight to the point.

These slides are based on the book
Supercomputing for Artificial Intelligence (Torres, 2025).
more info: <https://torres.ai/hpc4aibook/>

PDF slides are freely available for students.

Teachers using this book may request the PPTX version for
classroom use.