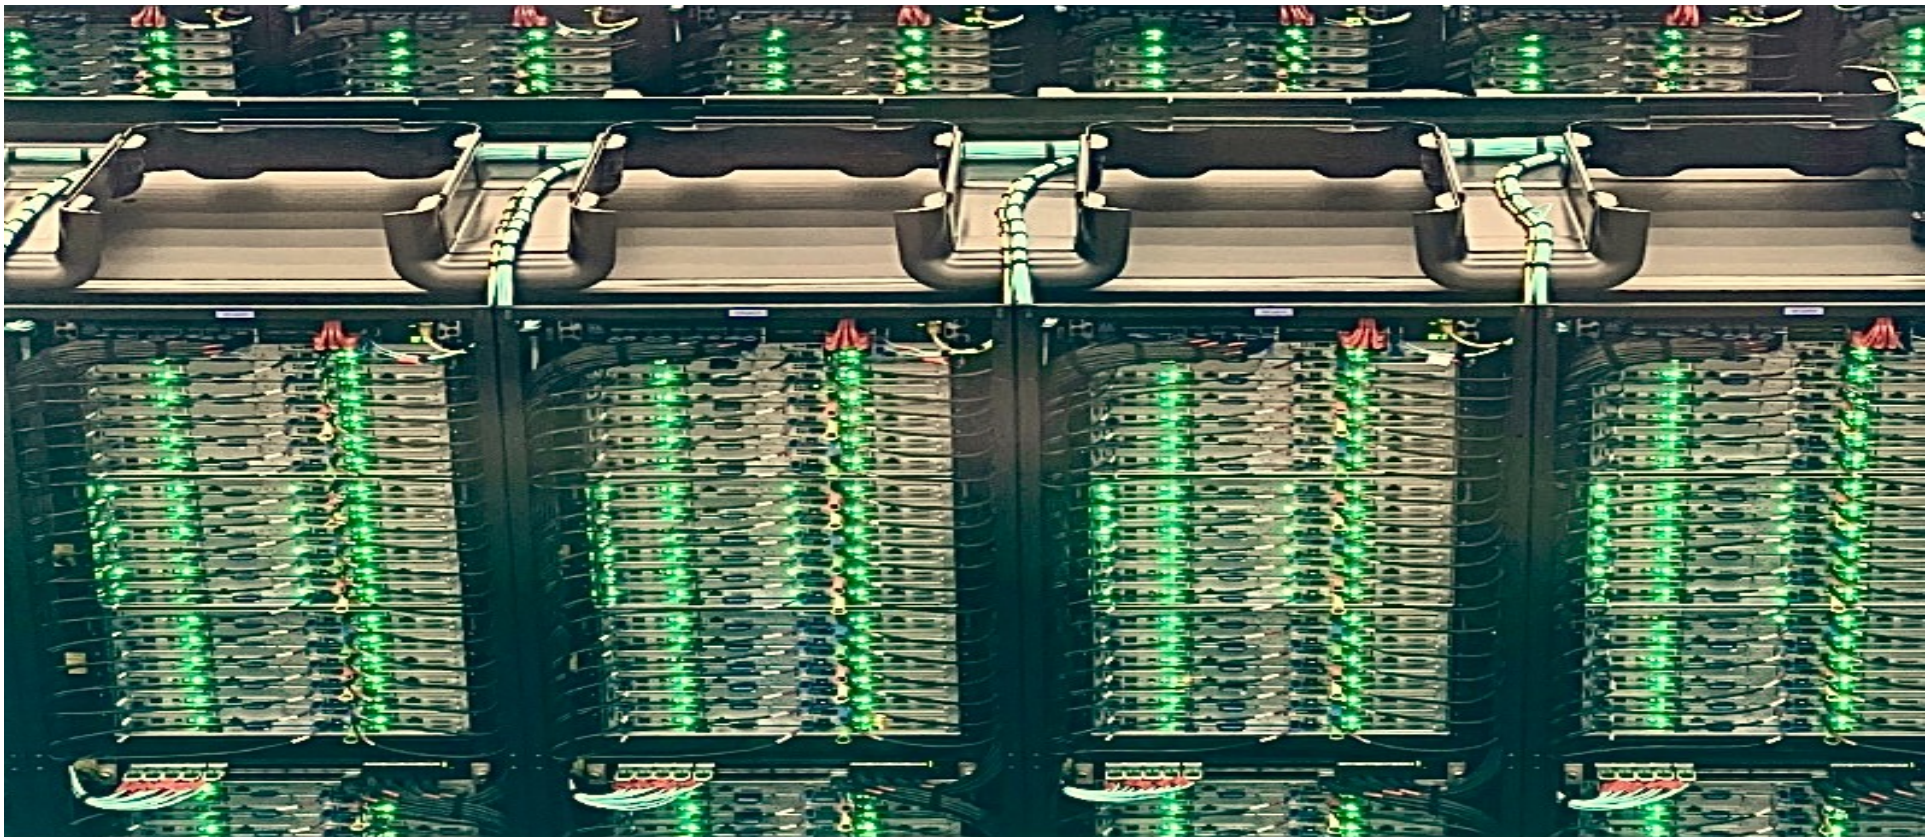


04. Launching and Structuring Parallel Programs with MPI

Supercomputing for Artificial Intelligence
Foundations, Architectures, and Scaling Deep Learning Workloads

Jordi **TORRES.AI**



Content

4.1 Foundations of Parallel Execution on Supercomputers

- How We Launch Parallel Jobs in This Book
- Why Parallel Programming Models Matter

4.2 Getting Started with MPI


- MPI Hello World
- mpirun Launcher
- srun Launcher
- Key MPI Concepts

4.3 Case Study: Parallelizing the Trapezoidal Rule Example

- Serial Implementation
- MPI Parallel Version
- Handling I/O in MPI Programs
- Barrier Synchronization

4.4 Collective Communication Primitives Broadcast

- Reduction Operations
- Scatter and Gather
- Advanced Patterns



4.1 Foundations of Parallel Execution on Supercomputers

Foundations of Parallel Execution

- **Running parallel programs requires:**
 - A mechanism to **launch jobs**
 - A model to **express parallelism**
- **This section:**
 - Introduces **srun** (used throughout this book)
 - Explains the motivation behind **traditional parallel programming models** (MPI, OpenMP)
- **Goal: understand how workloads execute efficiently at scale**

Launching Parallel Jobs

- Once SLURM allocates resources, it does not control execution
- Users must decide how to launch parallel work
- Two main approaches:
 - mpirun (traditional MPI launcher)
 - Still widely used in HPC to launch MPI applications
 - Works without workload managers like SLURM
 - Starts multiple program instances across nodes
 - mpiexec is usually synonymous with mpirun
 - srun (SLURM-integrated launcher)

srun

- **Native SLURM launcher, tightly integrated with resource management**
- **Syntax:** `srun [options] <executable>`
- **Options:**
 - `-n, --ntasks` → number of tasks
 - `-c, --cpus-per-task` → CPU cores per task
 - `-N, --nodes` → number of nodes
- **Example:**

```
srun -n4 -c8 ./my_app
```
- **Inherits parameters from #SBATCH directives**

Parallel Job Strategies in this Book

- **Default: srun for most examples on MareNostrum 5**
- **Exceptions:**
 - **mpirun** examples → for learning
 - **TensorFlow (MirroredStrategy)** → no srun, handled internally
- **All frameworks benefit from Singularity containers**
- **Next Table summarizes launch strategies by framework & chapter**

Parallel Job Strategies in this Book

Programming Model / Framework	Parallelization Scope	Launch Strategy Adopted in This Book	Book Chapter
MPI	multi-node	<code>srun -n N ./my-mpi-app</code> <code>mpirun -np N ./my-mpi-app</code>	Chapter 4
CUDA	single-GPU	<code>srun ./my-cuda-app</code>	Chapter 5
MPI + CUDA	multi-node multi-GPU	<code>mpirun -np N ./cuda-mpi-app</code>	Chapter 6
TensorFlow (MirroredStrategy)	single-node multi-GPU	<code>singularity exec python script.py</code>	Chapter 10
PyTorch (torchrun ²⁰)	multi-node, multi-GPU	<code>srun singularity exec torchrun ...</code>	Chapter 11 Chapter 12 Chapter 15

Why Parallel Programming Models Matter

- Launching jobs is not enough
 - need to define *how work is divided*
- Key historical idea: divide to conquer
- Early HPC systems → required abstractions for parallelism
 - **MPI** for distributed-memory systems
 - **OpenMP** for shared-memory systems
- These models hide low-level hardware details while enabling scalability

MPI, NCCL, and OpenMP

- **MPI (1990s)**
 - 1990s
 - Portable, scalable, de facto standard in scientific HPC
 - Still vital in some scientific domains
- **Modern systems use CPUs + GPUs**
 - New models required
- **NCCL**
 - Optimized for GPU collectives, now default in PyTorch
- **OpenMP**
 - Still important for CPU-based shared-memory workloads, but not used in large-scale AI training
- **In this book:**
 - Focus on MPI + CUDA + NCCL hybrid approaches



4.2 Getting Started with MPI

What is MPI?

- **Message Passing Interface (MPI)**
 - Standardized, portable library
 - Not a language → callable from C, C++, Fortran
 - Programmer must manage explicit communication
 - Enables processes to cooperate via explicit send/receive
 - De facto standard for distributed-memory HPC programming
 - Each process has its own address space
 - Enables growth by adding more nodes
 - Portable across supercomputers worldwide

Distributed Memory in MPI

- Each process has its own address space
- Excellent scalability: each node brings CPU, memory, storage, bandwidth
- Enables growth by adding more nodes
- Cost: programmer must manage explicit communication
- Shared data structures are harder to map
 - use message passing is required

MPI Hello World

- Each process prints:
 - Its rank (unique ID)
 - The total number of processes

```
#include <stdio.h>
#include <mpi.h>

int main (int argc, char **argv) {
    int rank, size;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    printf("I am %d of %d\n", rank, size);

    MPI_Finalize();
    return 0;
}
```

Compiling with Intel MPI

- **Intel wrapper compiler: mpiicx**
 - Combines Intel icx with MPI flags & libraries
- **Steps:**
 - Load Intel oneAPI MPI module

```
module load oneapi/2023.2.0
```

- Compile program

```
mpiicx mpi_helloworld.c -o mpi_helloworld
```

Launching with mpirun

- **mpirun** is the “classic” MPI launcher
- Starts multiple program instances, one per process
- **Handles:**
 - Process spawning
 - Rank assignment
 - Communication setup

- **Example:**

```
mpirun -np 4 ./mpi_helloworld
```

→ Output (unordered):

```
I am 2 of 4  
I am 3 of 4  
I am 0 of 4  
I am 1 of 4
```

Launching with srun

- Native SLURM launcher: **srun**
- Integrated with resource allocation & accounting
- **mpirun** and **srun** launchers work, but **srun** is preferred on MN5
- Example batch script runs program twice:
 - First with **mpirun**
 - Then with **srun**

```
echo "mpirun ./mpi_helloworld:"  
mpirun ./mpi_helloworld  
  
echo "srun ./mpi_helloworld"  
srun ./mpi_helloworld
```

Launching MPI in SLURM

```
#!/bin/bash
#SBATCH -J mpi_helloworld
#SBATCH -t 00:15
#SBATCH -o %x_%J.out
#SBATCH -e %x_%J.err
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=1
#SBATCH --exclusive
#SBATCH --account <account>
#SBATCH --qos gp_debug

module load oneapi/2023.2.0

echo "mpirun ./mpi_helloworld:"
mpirun ./mpi_helloworld

echo "srun ./mpi_helloworld"
srun ./mpi_helloworld
```

– Example:

- SLURM directive `--ntasks=4` ensures 4 processes are launched
- Both `mpirun` and `srun` launch 4 processes
- Each prints its rank (0...3)
- Order is not guaranteed → asynchronous output
- **From now on:**
use `srun` in SLURM-managed environments

Launching MPI in SLURM

```
mpirun ./mpi_helloworld:
```

```
I am 2 of 4
```

```
I am 1 of 4
```

```
I am 3 of 4
```

```
I am 0 of 4
```

```
srun ./mpi_helloworld
```

```
I am 3 of 4
```

```
I am 0 of 4
```

```
I am 2 of 4
```

```
I am 1 of 4
```

Task 4.1

- **Task 4.1 – Compile and run your first MPI program**
 - a) Compile and execute the basic MPI Hello World example using the `mpicx` compiler.
 - b) Use the SLURM script provided to launch 4 MPI tasks.
 - c) Examine the output and confirm that each process prints its rank and the total number of processes.

Multi-node Distribution Example

- Goal: run 16 tasks across 4 nodes (4 per node)
- SLURM directives:

```
#SBATCH --ntasks=16          # Total number of MPI tasks
#SBATCH --nodes=4             # Number of nodes to allocate
#SBATCH --ntasks-per-node=4   # Number of tasks per node
```

Multi-node Distribution Example

- C code prints process rank + hostname

```
#include <stdio.h>
#include <mpi.h>
#include <unistd.h> // for gethostname

int main (int argc, char **argv) {
    int rank, size;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    char hostname[256];
    gethostname(hostname, sizeof(hostname));
    printf("I am %d of %d running on %s\n", rank, size, hostname);

    MPI_Finalize();
    return 0;
}
```

Multi-node Distribution Example

- **Compile the program:**

```
module load oneapi/2023.2.0
mpiicx -o mpi_helloworld4x4 mpi_helloworld4x4.c
```

- **SLURM batch script:**

```
#!/bin/bash
#SBATCH -J mpi_helloworld_4x4
#SBATCH -t 00:15
#SBATCH -o %x_%J.out
#SBATCH -e %x_%J.err
#SBATCH --ntasks=16
#SBATCH --nodes=4
```

```
#SBATCH --ntasks-per-node=4
#SBATCH --cpus-per-task=1
#SBATCH --exclusive
#SBATCH --account <account>
#SBATCH --qos gp_debug

srun ./mpi_helloworld4x4
```


Multi-node Distribution Example

```
I am 1 of 16 running on gs26r2b52
I am 2 of 16 running on gs26r2b52
I am 3 of 16 running on gs26r2b52
I am 0 of 16 running on gs26r2b52
I am 12 of 16 running on gs26r2b72
I am 4 of 16 running on gs26r2b62
I am 8 of 16 running on gs26r2b70
I am 5 of 16 running on gs26r2b62
I am 13 of 16 running on gs26r2b72
I am 6 of 16 running on gs26r2b62
I am 14 of 16 running on gs26r2b72
I am 7 of 16 running on gs26r2b62
I am 15 of 16 running on gs26r2b72
I am 9 of 16 running on gs26r2b70
I am 10 of 16 running on gs26r2b70
I am 11 of 16 running on gs26r2b70
```

- Tasks are spread across nodes:
 - Output shows ranks mapped to hostnames (gs26r2b52, ...)

Task 4.2

■ Task 4.2 – Observe node distribution using hostnames

- a) Modify the Hello World program to print the hostname (use `gethostname()`).
- b) Then, rerun the job with `--ntasks=16` and `--nodes=2` to distribute 16 MPI tasks over 2 nodes.
- c) Use `--ntasks-per-node=8` to enforce the distribution.
- d) Verify by analyzing the standard output that the results are correct and that the execution took place across two different nodes, as expected.
- e) Try running the program using both `srun` and `mpirun`.

Key MPI Concepts

- **SPMD model: Single Program, Multiple Data**
 - Same binary → different behavior per rank
- **Every program must:**
 - MPI_Init() → enter MPI world
 - MPI_Finalize() → clean up MPI resources
- **When a parallel MPI program starts, all participating processes are grouped into a default communicator (called MPI_COMM_WORLD).**

Communicators and Ranks

- `MPI_Comm_rank(MPI_COMM_WORLD, &rank)`
 - Returns rank (0 ... size-1)
- `MPI_Comm_size(MPI_COMM_WORLD, &size)`
 - Returns total processes
- **Enables:**
 - Work distribution
 - Conditional logic (rank 0 often handles I/O, coordination, ...)
- **IMPORTANT: Each process gets:**
 - **Rank:** unique ID
 - **Size:** total number of processes

Point-to-Point Communication

- **Core MPI operations:**

- `MPI_Send(...)` → send data
- `MPI_Recv(...)` → receive data

- **Requirements:**

- Same communicator
- Matching source/destination ranks
- Matching tags and message size/type
- If `MPI_Recv` has no matching send
 - blocks indefinitely

Example: Sending Greetings

- Non-zero processes send greetings to rank 0

```
#include <stdio.h>
#include <string.h>
#include <mpi.h>

int main(int argc, char* argv[]) {
    int rank, size;
    char message[100];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (rank != 0) {
        // All processes except 0 send a message to process 0
        sprintf(message, "Greetings from process %d!", rank);
        MPI_Send(message, strlen(message) + 1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
    } else {
        // Process 0 receives messages from all other processes
        for (int source = 1; source < size; source++) {
            MPI_Recv(message, 100, MPI_CHAR, source, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("Process 0 received: %s\n", message);
        }
    }

    MPI_Finalize();
    return 0;
}
```

Example: Sending Greetings

- Rank 0 loops, receives messages sequentially
- Uses blocking `MPI_Recv`
 - ordered reception by rank

```
if (rank != 0) {  
    // All processes except 0 send a message to process 0  
    sprintf(message, "Greetings from process %d!", rank);  
    MPI_Send(message, strlen(message) + 1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);  
} else {  
    // Process 0 receives messages from all other processes  
    for (int source = 1; source < size; source++) {  
        MPI_Recv(message, 100, MPI_CHAR, source, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
        printf("Process 0 received: %s\n", message);  
    }  
}
```

Example: Sending Greetings

- Sample output with 8 tasks:

```
Process 0 received: Greetings from process 1!  
Process 0 received: Greetings from process 2!  
Process 0 received: Greetings from process 3!  
Process 0 received: Greetings from process 4!  
Process 0 received: Greetings from process 5!  
Process 0 received: Greetings from process 6!  
Process 0 received: Greetings from process 7!
```

Task 4.3

■ Task 4.3 – Point-to-point communication

- a) Compile and run the MPI program using `MPI_Send` and `MPI_Recv` where all worker processes send a greeting to process 0.
- b) Create 16 tasks.
- c) Confirm that process 0 receives the messages in rank order.




4.3

Case Study: Parallelizing the Trapezoidal Rule

Case Study: Trapezoidal Rule (MPI)

- **Goal:**
 - approximate a definite integral with the trapezoidal rule
- **Strategy:**
 - split $[a, b]$ into n equal subintervals; sum trapezoids
- **Parallel idea:**
 - each process sums a subset of trapezoids, then reduce

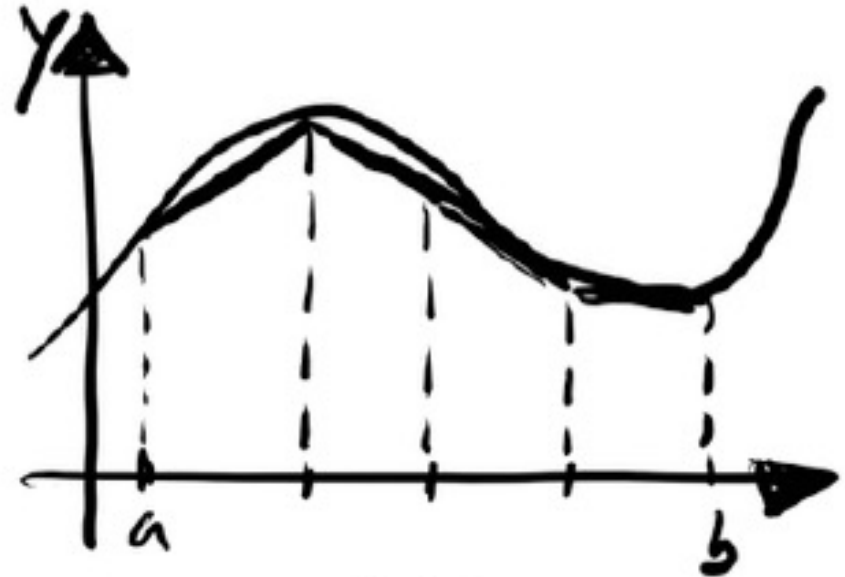
$$\int_a^b f(x) dx$$



Trapezoidal Rule

■ Pseudo-code:

```
/* Input: a, b, n */  
h = (b - a) / n;  
approx = (f(a) + f(b)) / 2.0;  
for i = 0 .. n-1:  
    x_i = a + i*h  
    approx += f(x_i)  
approx *= h
```



■ Accuracy ↑ with n

Serial Example: π

- We approximate:

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$

- C implementation:

```
double f(double x){ return 4.0/(1.0+x*x); }
```

```
double trap(double a,double b,int n){  
    double h=(b-a)/n, sum=(f(a)+f(b))/2.0;  
    for(int i=1;i<n;i++) sum += f(a + i*h);  
    return sum*h;  
}
```


Serial Example: π

```
double f(double x) {
    return 4.0 / (1.0 + x * x);
}

double trapezoidal_rule(double a, double b, int n) {
    double h = (b - a) / n;
    double sum = (f(a) + f(b)) / 2.0;

    for (int i = 1; i < n; i++) {
        sum += f(a + i * h);
    }

    return sum * h;
}

int main() {
    int n = 1000000;
    double a = 0.0, b = 1.0;
    double pi = trapezoidal_rule(a, b, n);
    printf("Estimated PI = %.16f\n", pi);
    return 0;
}
```

Serial Example: π

- **Compile and Run:**

```
$ module load intel  
$ icx -O3 pi_seq.c -o pi_seq  
$ ./pi_seq
```

Estimated PI = 3.1415926535895844


Task 4.4

■ Task 4.4 – Write and Run The Sequential Program That Estimate π

Use a standard C compiler and verify that the estimated value of π is printed to the terminal. This simple test ensures that your code is working before proceeding to the parallel MPI version.

Parallelization Plan

- **Partition work across p processes:**
 - Each process computes n/p trapezoids over its subinterval
- **Communication pattern:**
 - Local compute \rightarrow send partial sums \rightarrow root does final sum
- **Parallel reduction (manual or MPI_Reduce later)**



MPI Parallel Implementation

```
#include <stdio.h>
#include <mpi.h>

double f(double x) {
    return 4.0 / (1.0 + x * x);
}

double local_trap(double a, double b, int local_n) {
    double h = (b - a) / local_n;
    double sum = (f(a) + f(b)) / 2.0;

    for (int i = 1; i < local_n; i++) {
        sum += f(a + i * h);
    }

    return sum * h;
}
```

MPI Parallel Implementation

```
int main(int argc, char** argv) {  
    int rank, size, n;  
    double a = 0.0, b = 1.0;  
    double local_a, local_b, h;  
    int local_n;  
    double local_result, total_result;  
  
    MPI_Init(&argc, &argv);  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);
```

MPI Parallel Implementation

```
// Ensuring that n is divisible by the number of MPI processes
n = 1000000;
n = (n / size) * size;

h = (b - a) / n;
local_n = n / size;
local_a = a + rank * local_n * h;
local_b = local_a + local_n * h;

local_result = local_trap(local_a, local_b, local_n);
```

MPI Parallel Implementation

```
if (rank == 0) {
    total_result = local_result;
    for (int source = 1; source < size; source++) {
        double temp;
        MPI_Recv(&temp, 1, MPI_DOUBLE, source, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        total_result += temp;
    }
    printf("Estimated PI = %.16f\n", total_result);
} else {
    MPI_Send(&local_result, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
}

MPI_Finalize();
return 0;
}
```


Compile & Run (MPI)

■ Compile

```
module load oneapi/2023.2.0  
mpiicx pi_mpi.c -o pi_mpi
```

■ SLURM script (CPU only):

```
#!/bin/bash  
#SBATCH --job-name=pi_mpi  
#SBATCH --output=%x_%j.out  
#SBATCH --error=%x_%j.err  
#SBATCH --ntasks=8  
#SBATCH --cpus-per-task=1  
#SBATCH --time=00:05:00  
#SBATCH --exclusive  
#SBATCH --account=<account>  
#SBATCH --qos=gp_debug  
  
srun ./pi_mpi
```

Task 4.5

■ Task 4.5 – Write and Run the Parallel MPI Code to Estimate the Value of π

Execute the parallel version using 16 processes with SLURM, and verify that the estimated value of π is correct, as this will serve as the basis for the following tasks.

Taking Timing

- **We need performance metrics**
 - measure execution time of code under study
- **Not just total program time; often a specific region**
- **Typical timers:**
 - MPI: `MPI_Wtime()`
 - OpenMP: `omp_get_wtime()`
 - **POSIX (portable):** `gettimeofday()`

POSIX gettimeofday()

[Appendix 17.4]

- **Portable Wall-Clock**
- **Works in serial, MPI, OpenMP, mixed codes**
- **Resolution: microseconds (10^{-6} s)**
- **Prototype:**

```
#include <sys/time.h>
int gettimeofday(struct timeval * tv, struct timezone * tz);
```

- **Structure:**

```
struct timeval {
    time_t tv_sec;      /* seconds */
    suseconds_t tv_usec; /* microseconds */
};

struct timezone {
    int tz_minuteswest; /* minutes west of Greenwich */
    int tz_dsttime;     /* type of DST correction */
};
```

Timing Pattern

■ (Copy/Paste Template)

```
#include <sys/time.h>
struct timeval start_time, end_time;

gettimeofday(&start_time, NULL);

// << code section to measure >>

gettimeofday(&end_time, NULL);
print_times();
}

print_times()
{
    int total_usecs;
    float total_time;
    total_usecs = (end_time.tv_sec - start_time.tv_sec) * 1000000 +
                  (end_time.tv_usec - start_time.tv_usec);
    printf(" %.2f mSec \n", ((float) total_usecs) / 1000.0);
    total_time = ((float) total_usecs) / 1000000.0;
}
```

Mini Example (fills a matrix)

```
/* timesample.c */
#include <stdio.h>
#include <sys/time.h>
#include <stdlib.h>

#define SIZE 1000
typedef double matrix[SIZE][SIZE];
matrix m1;

struct timeval start_time, end_time;

static void foo(void) {
    int i, j;
    for (i = 0; i < SIZE; ++i)
        for (j = 0; j < SIZE; ++j)
            m1[i][j] = 1.0;
}
```

Mini Example (fills a matrix)

```
void print_times() {  
    int total_usecs;  
    float total_time;  
    total_usecs = (end_time.tv_sec - start_time.tv_sec) * 1000000 +  
                  (end_time.tv_usec - start_time.tv_usec);  
    printf(" %.2f mSec \n", ((float) total_usecs) / 1000.0);  
    total_time = ((float) total_usecs) / 1000000.0;  
}
```

Mini Example (fills a matrix)

```
int main() {  
    int i;  
  
    gettimeofday(&start_time, NULL);  
  
    for (i = 0; i < 10; ++i) {  
        foo();  
    }  
  
    gettimeofday(&end_time, NULL);  
    print_times();  
  
    return 0;  
}
```


Compile & Run (example output)

```
$ gcc -O3 -o timesample timesample.c  
$ ./timesample
```

3.81 mSec

How does it scale with varying numbers of parallel processes?

■ Compile

```
module load oneapi/2023.2.0
```

```
icx -O3 pi_seq_timed.c -o pi_seq_timed
```

```
mpiicx -O3 pi_mpi_timed.c -o pi_mpi_timed
```

The files `pi_mpi_timed.c` and `pi_seq_timed.c` are the same as the previous versions, but now include the necessary code to measure execution time using the `gettimeofday` system call.

How does it scale with varying numbers of parallel processes?

- `pi_mpi_scaling.slurm`

```
#!/bin/bash
#SBATCH --job-name=pi_scaling
#SBATCH --output=%x_%j.out
#SBATCH --error=%x_%j.err
#SBATCH --ntasks=64
#SBATCH --cpus-per-task=1
#SBATCH --time=00:15:00
#SBATCH --exclusive
#SBATCH --account=<account>
#SBATCH --qos=gp_debug

./pi_seq_timed

for P in 2 4 8 16 32 64
do
    echo ""
    srun --ntasks=$P ./pi_mpi_timed
done
```

How does it scale with varying numbers of parallel processes?

Processes	Time (ms)	Speedup	Efficiency (%)
seq	4430.33	1.00	100.0%
2	2394.17	1.85	92.4%
4	1198.30	3.70	92.6%
8	601.04	7.37	92.1%
16	308.02	14.38	89.9%
32	160.61	27.58	86.2%
64	216.27	20.49	32.0%

How does it scale with varying numbers of parallel processes?

- We can observe that up to 32 processes, the efficiency remains reasonable; however, beyond that point, it drops significantly.
- In fact, the execution time with 64 processes is even higher than with 32, indicating diminishing returns.
- This raises the question: is the program's limited parallelism the cause?

→ See Amdahl's Law & Gustafson's Law sections in the book

Task 4.6 (from book)

■ ~~Task 4.6 – Analysis using Gustafson’s Law to estimate the value of π~~

~~Run the parallel version of the program that estimates the value of π using up to 128 processes, and analyze the results according to Gustafson’s Law, following the same methodology presented earlier in this chapter.~~

Task 4.6 (new)

■ Task 4.6 new – Scalability Analysis of the MPI Trapezoidal Rule

- a) Compile the provided sequential and parallel timed versions of the trapezoidal rule program (pi_seq_timed.c and pi_mpi_timed.c).
- b) Submit the SLURM script pi_mpi_scaling.slurm, which runs the parallel program with different numbers of processes (2, 4, 8, 16, 32) under the same job allocation.
- c) Compare execution times, calculate speedup and efficiency, and discuss at which point adding more processes no longer improves performance.
- d) Reflect on possible causes such as communication overhead or limited parallelism.

MPI I/O Basics

- **stdout from many ranks can interleave (unordered prints)**
- **stdin typically only from rank 0**
 - Distribute input to others:

```
void get_input(int* n_p, int my_rank, int comm_sz) {  
    if (my_rank == 0) {  
        printf("Enter number of intervals: ");  
        scanf("%d", n_p);  
  
        for (int dest = 1; dest < comm_sz; dest++) {  
            MPI_Send(n_p, 1, MPI_INT, dest, 0, MPI_COMM_WORLD);  
        }  
    } else {  
        MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
    }  
}
```

- Later: simpler with **MPI_Bcast**

Barrier Synchronization

- **Ensure all ranks reach a point before proceeding:**

```
int MPI_Barrier(MPI_Comm comm);
```

- **Use around timed regions to align start/stop**
- **Caveat: overuse introduces idle waiting → hurt perf**



4.4 Collective Communication Primitives

Collective Communication in MPI

- **Point-to-point operations (MPI_Send, MPI_Recv) are flexible but ...**
- **Many parallel algorithms need global communication patterns**
- **MPI provides collective operations:**
 - Simplify programming and code maintenance
 - Highly optimized in most implementations
- **All processes in the communicator must call them, otherwise deadlock**

MPI_Bcast (Broadcast)

- Used when one process (root) sends the same data to all others
- Function prototype:

```
MPI_Reduce(&local_result, &total_result, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

- Example use case: input read by process 0 distributed to all others
- More elegant and efficient than multiple MPI_Send / MPI_Recv

Reduction Operations

- **Goal: combine partial results from all processes**
- **Traditional approach: manual send/receive loop**
- **Collective version:**

```
MPI_Reduce(&local_result, &total_result, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

- **Variants:**
 - **MPI_Reduce** → result only at root
 - **MPI_Allreduce** → result delivered to all processes
- **Operations supported: MPI_SUM, MPI_MAX, MPI_MIN, ...**

MPI_Scatter and MPI_Gather

- **MPI_Scatter**
 - root splits dataset into chunks → each process gets one
- **MPI_Gather**
 - each process sends results → root collects into array
- **Typical workflow:**
 - Root distributes a vector (`MPI_Scatter`)
 - Processes compute locally
 - Root collects updated vector (`MPI_Gather`)
- **Example: parallel vector sum** (available in repository)

Task 4.7

■ Task 4.7 – Experimenting with Scatter and Gather

- a) Take the example provided in the book's GitHub repository and run it.
- b) Once you have verified that it executes correctly, modify the code to test different vector sizes.
- c) Analyze how the size of the vector influences the program's speedup, efficiency, and scalability.

Advanced Collectives (I)

■ **MPI_Allgather**

- Every process gathers contributions from all others
- Equivalent to Gather + Broadcast, but more efficient
- Example: sharing local statistics with all processes

■ **MPI_Alltoall**

- Each process sends/receives different data to/from all others
- General redistribution primitive, used in sorting, domain decomposition, matrix transpose

Advanced Collectives (II)

- **MPI_Reduce_scatter**
 - Combines reduction and scatter in a single step
 - Each process receives only its needed portion of the reduced result
 - Useful in distributed dot products, localized data aggregation
- **Advantages of advanced collectives:**
 - Reduce code complexity
 - Leverage optimized implementations for scalability

Pc: MPI

■ **Tasks included:**

Task 4.1

Task 4.2

Task 4.3

Task 4.4

Task 4.5

Task 4.6 (new)

Task 4.7

■ **Deliverable:**

- Upload a single PDF (per group) to the intranet racó@FIB containing one slide per task. Each slide should report results or briefly explain how the task was completed.
- In class (evaluation day), one group (chosen at random) will give a “*elevator pitch*”-style or PK presentation — clear, concise, and straight to the point.

These slides are based on the book
Supercomputing for Artificial Intelligence (Torres, 2025).
more info: <https://torres.ai/hpc4aibook/>

PDF slides are freely available for students.

Teachers using this book may request the PPTX version for
classroom use.