

# SA-MIRI Course Schedule (22/10 Tentative)

## DL basics Week

dav	Wednesday	dav	Friday
22/10	7. Neural Networks: Concepts and First Steps Pf part 1	Presentation Pa Presentation Pb	24/10 8. Training Neural Networks: Basics, CNNs, and Deployment 9. Getting Started with PyTorch Pf part 2
29/10	10. Introduction to Parallel Training of Neural Networks	Presentation Pc Presentation Pd Presentation Pe	31/10 Midterm period (no class)
05/11	(Midterm period) PRESENTATIONS—I Midterm period (no class)		07/11 Pg
12/11	11. Practical Guide to Efficient Training with PyTorch 12. Parallelizing Model Training with Distributed Data Parallel	Presentation Pf	14/11 Ph
19/11	Ph	Presentation Pg	21/11 Pi
26/11	Honoris Causa Dr. Oriol Vinyals (no class)		28/11 Pi
03/12	13. Introduction to Large Language Models 14. End-to-End Large Language Models Workflow	Presentation Ph	05/12 Pj
10/12	15. Exploring Optimization and Scaling of LLMs	Presentation Pi	12/12 Pk*
17/12	Pk* 16. Looking Forward: Supercomputing and AI Futures	Presentation Pj	19/12 Pl * “attendance not required (for students returning home)”

\* No presentation, submission only.

# SA-MIRI Course Schedule (Tentative)

Practical Session	acronym	Book Chapter included	Practical Session	acronym	Book Chapter included
Pa	GETTING STARTED	task 2.1 – Log into MareNostrum 5 task 2.2 – Change your password task 2.3 – (Optional) Enable passwordless ssh authentication task 2.4 – Transfer files using scp task 2.5 – (Optional) Mount the MN5 filesystem on your laptop  task 3.1 – Compare icx and gcc compiler optimizations task 3.2 – Reflecting on slurm job prioritization task 3.3 – Submit your first slurm job	Pg	SCALING TF	task 10.1 – Task setup and file structure task 10.2 – Code review and understanding task 10.3 – Training on cpu (baseline performance) task 10.4 – Gpu execution time and cpu vs gpu comparison task 10.6 – Analyze the impact of gpu parallelism task 10.7 – Parallelization of resnet152 task 10.8 – Comparing parallel training performance across model sizes task 10.9 – Resnet101v2 scalability analysis  task 10.10 – Interpreting results with Amdahl's and Gustafson's laws task 10.11 – Applying Amdahl's and Gustafson's laws to Resnet101v2
Pb	CONTAINERS	task 3.4 – Install docker in your platform task 3.5 – Download docker image task 3.6 – Run docker image  task 3.7 – Stop a docker container task 3.8 – Run docker with port mapping task 3.9 – Start the jupyter notebook server task 3.10 – Create and run a test notebook	Ph	OPTIMIZE PT	task 11.1 – Find the maximum viable batch size task 11.2 – Investigating dataloader bottleneck with a lightweight model task 11.3 – Optimizing dataloaders with appropriate num_workers task 11.4 – Confirming dataloader efficiency with vit  task 11.5 – Enable mixed precision with vit + micro-224 task 11.6 – Reproducing the effect of torch.compile() task 11.7 – Report your conclusions task 11.8 – Minor code tweaks for a final throughput boost
Pc	MPI	task 4.1 – Compile and run your first mpi program task 4.2 – Observe node distribution using hostnames task 4.3 – Point-to-point communication task 4.4 – Write and run the sequential program that estimate $\pi$  task 4.5 – Write and run the parallel mpi code to estimate the value of $\pi$ task 4.6 – Analysis using Gustafson's law to estimate the value of $\pi$ task 4.7 – Experimenting with scatter and gather	Pi	SCALING PT	task 12.1 – Reproducing distributed training results on mn5 task 12.2 – Analyze and compare scaling efficiency task 12.3 – Investigate diminishing returns in training time task 12.4 – Find the sweet spot for your use case
Pd	CUDA	task 5.1 – Your first hello world in cuda task 5.2 – Dimensionality of a thread block and grid  task 5.3 – Investigating parallel execution with multiple threads task 5.4 – Element-wise vector addition using cuda task 5.5 – Parallel matrix multiplication with cuda  task 5.6 – Running cuda jobs with slurm task 5.7 – Profiling matrix multiplication on the gpu task 5.8 – Compute-bound vs memory-bound	Pj	END-TO-END HF	task 14.1 – Obtain your hugging face access token task 14.2 – Download and run the hello world in google colab task 14.3 – Download the model and dataset locally using huggingface-cli task 14.4 – Transfer the model and dataset to MareNostrum 5 task 14.5 – Run the inference and fine-tuning script on mn5  task 14.6 – Compare execution in colab vs mn5
Pe	CUDA-aware MPI	task 6.1 – Reflecting on cuda's execution model task 6.2 – Precision trade-offs: true or false? task 6.3 – Submit and validate the first performance run task 6.4 – Understand the metrics collection task 6.5 – Explore the effect of optimized compilation flags task 6.6 – Evaluate the impact of the cub library task 6.7 – Benchmarking the impact of gpu count and problem size	Pk	OPT & SCAL LLM	task 15.1 – Baseline experiment using facebook/opt-1.3b task 15.2 – Finding the out-of-memory limit task 15.3 – Mixed precision training task 15.4 – Model precision  task 15.5 – Increasing batch size with model precision task 15.6 – Enabling flash attention task 15.7 – Increasing batch size with flash attention task 15.8 – Using the liger kernel task 15.9 – Augmenting batch size due to liger kernels task 15.10 – Scaling on multiple gpus task 15.11 – Final reflections
Pf	DL COLABS	task 7.1 – Set up your google colab environment task 7.2 – Execute the provided notebook step-by-step task 7.3 – Improve the accuracy of your model  task 8.1 – Improving a basic cnn model task 8.2 – Exporting the python script task 8.3 – Running your first natural network on a login node task 8.4 – Submitting your first gpu dl training with slurm  task 9.1 – Comparative implementation in pytorch and tensorflow	Pl	FUTURE	task 16.1 – Mapping the pendulum shifts of the triad task 16.2 – Exploring the new giants of compute task 16.3 – Powering the future of ai task 16.4 – Charting the role of quantum in future supercomputing task 16.5 – From data farms to ai preserves task 16.6 – Embodiment as a source of learning

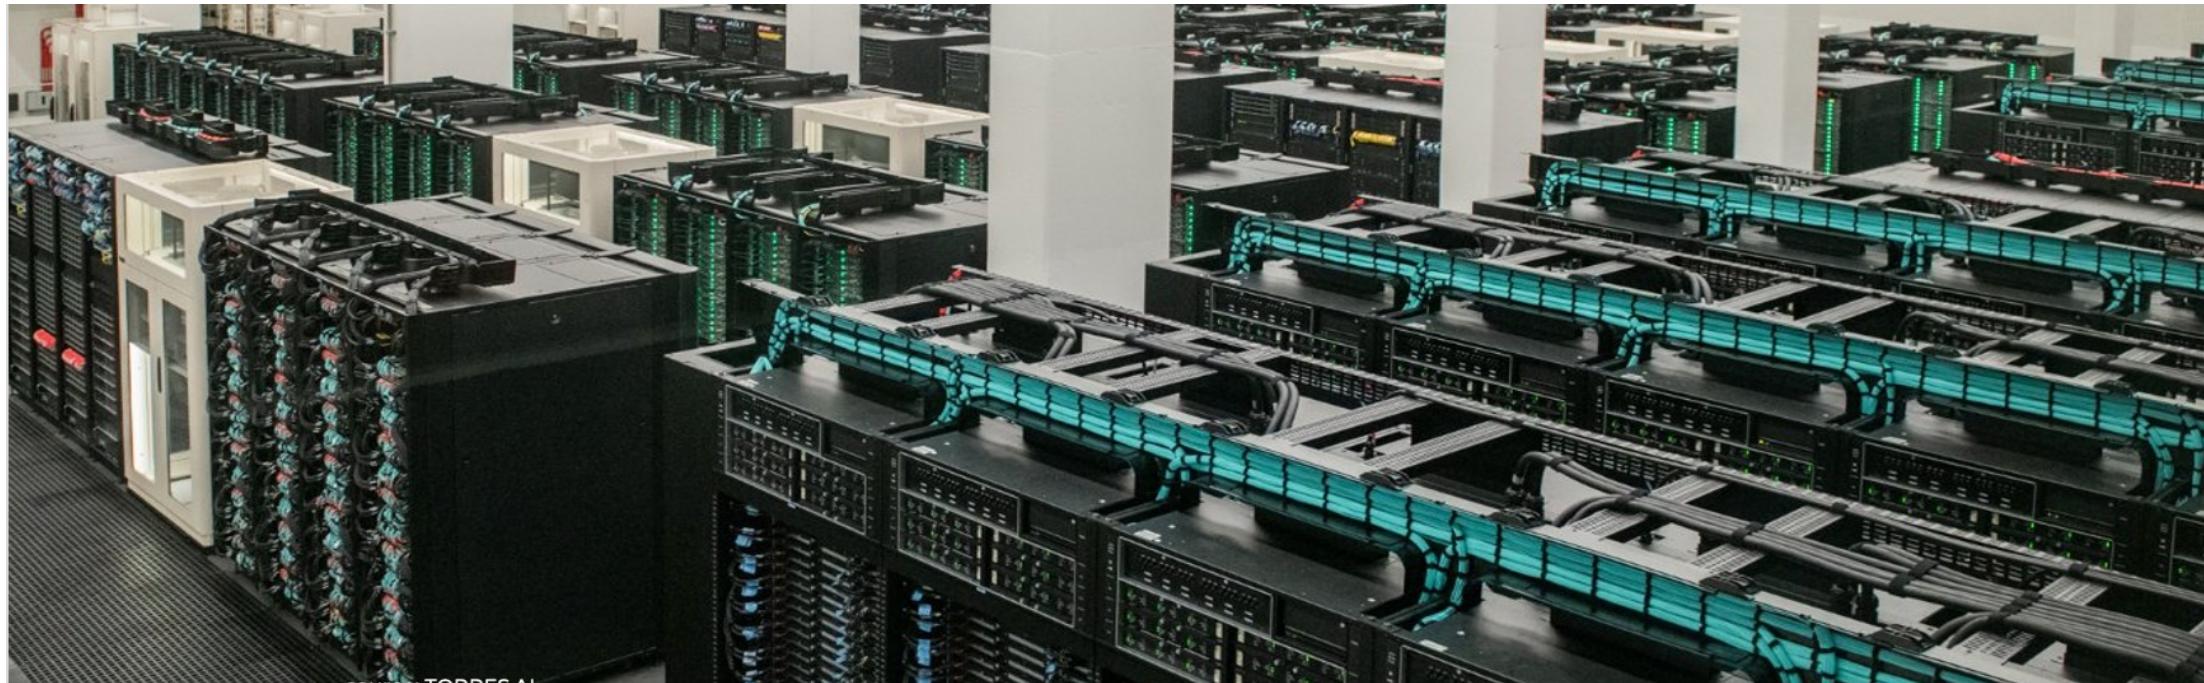
# GETTING STARTED WITH DEEP LEARNING (a recap combining Chapters 7, 8, and 9)

---

SUPERCOMPUTERS ARCHITECTURE

Master in Innovation and Research in Informatics

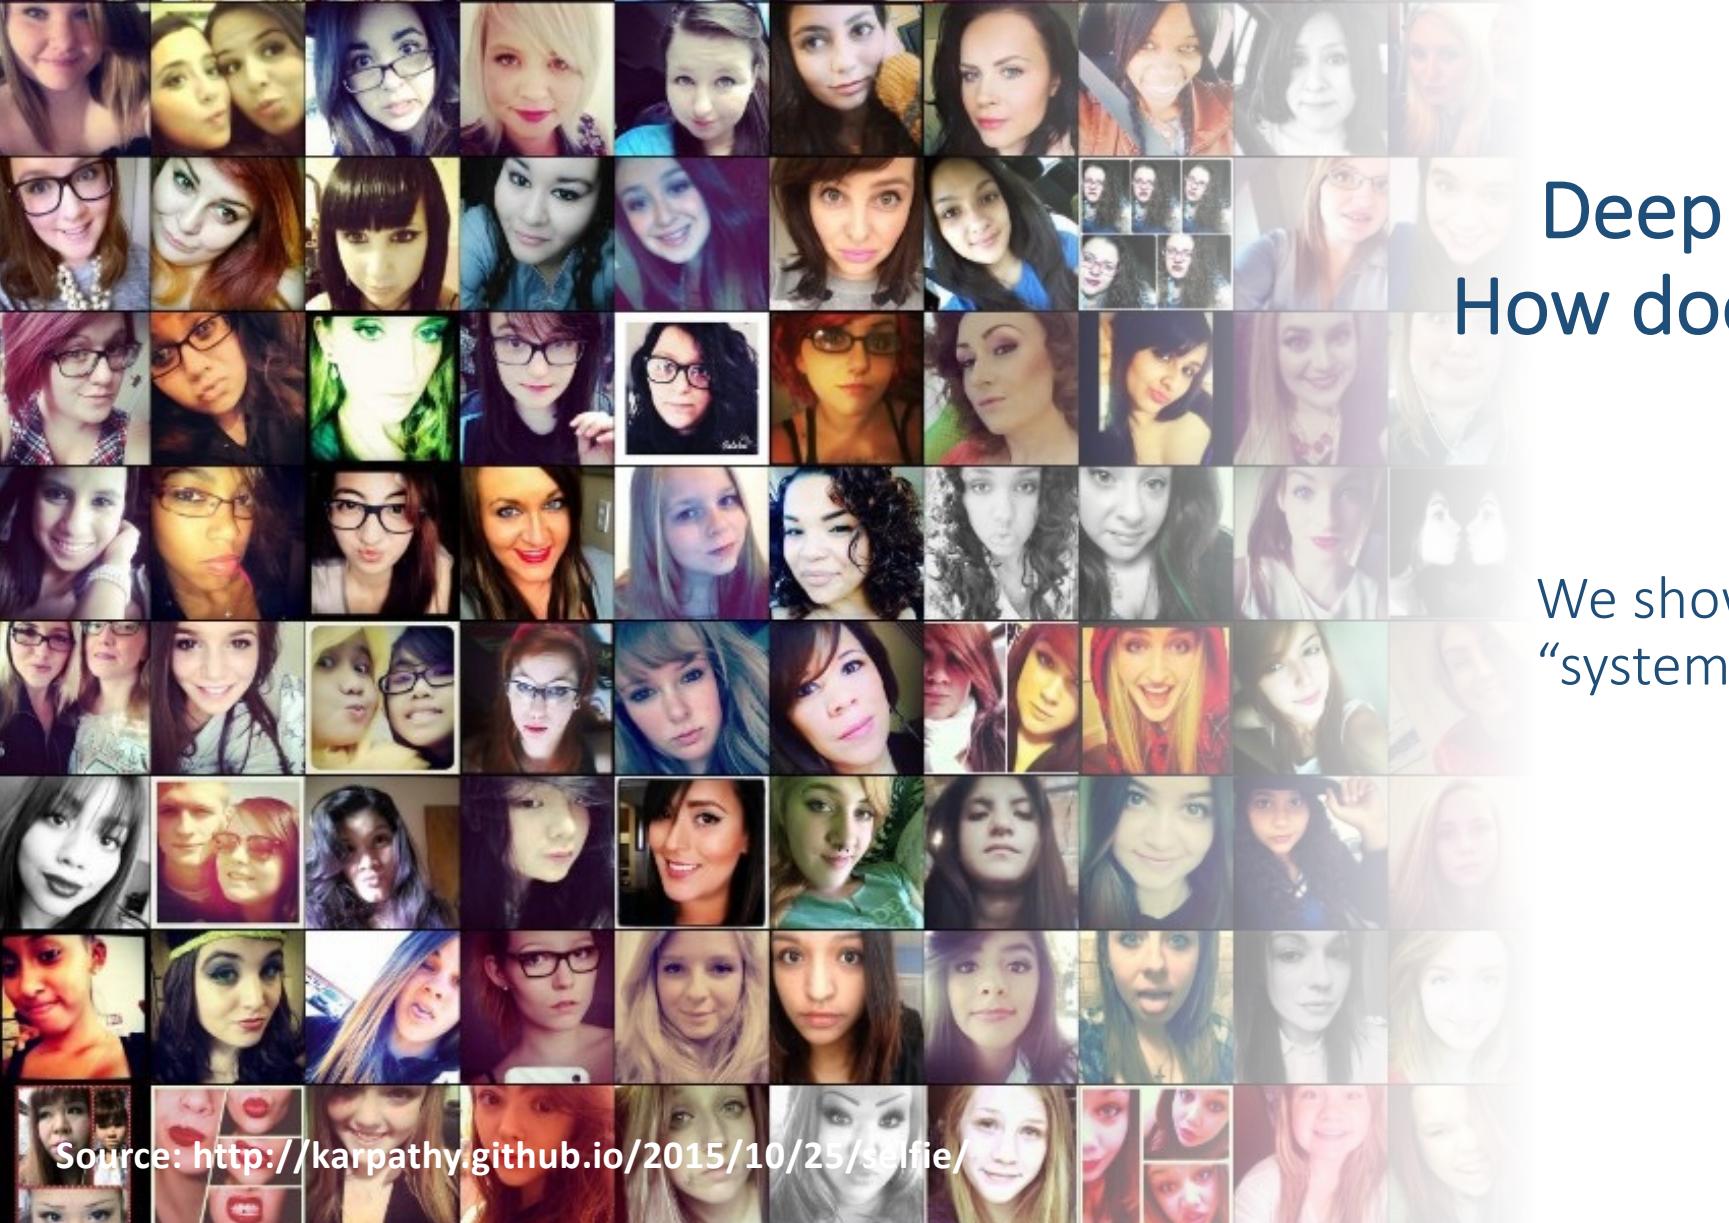
Jordi TORRES.AI





# How does Deep Learning Works?

Jordi **TORRES.AI**



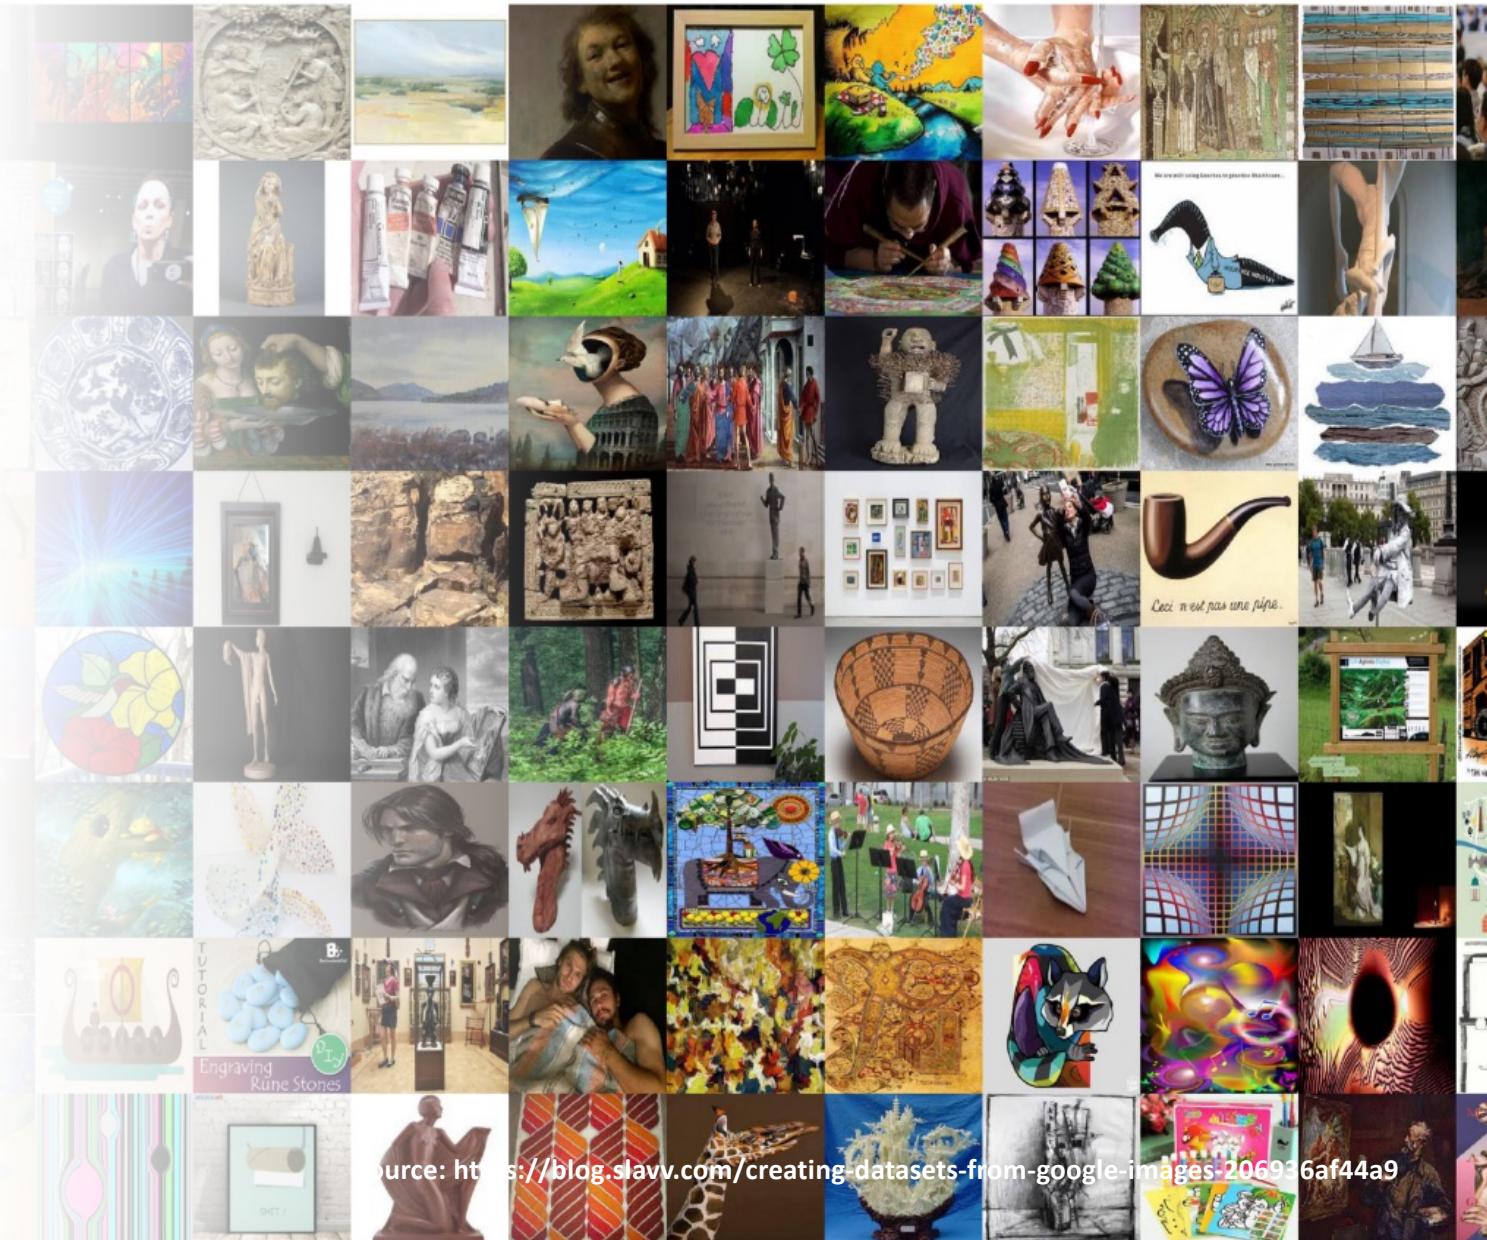
# Deep Learning How does it works?

We show to the  
“system” many FACES

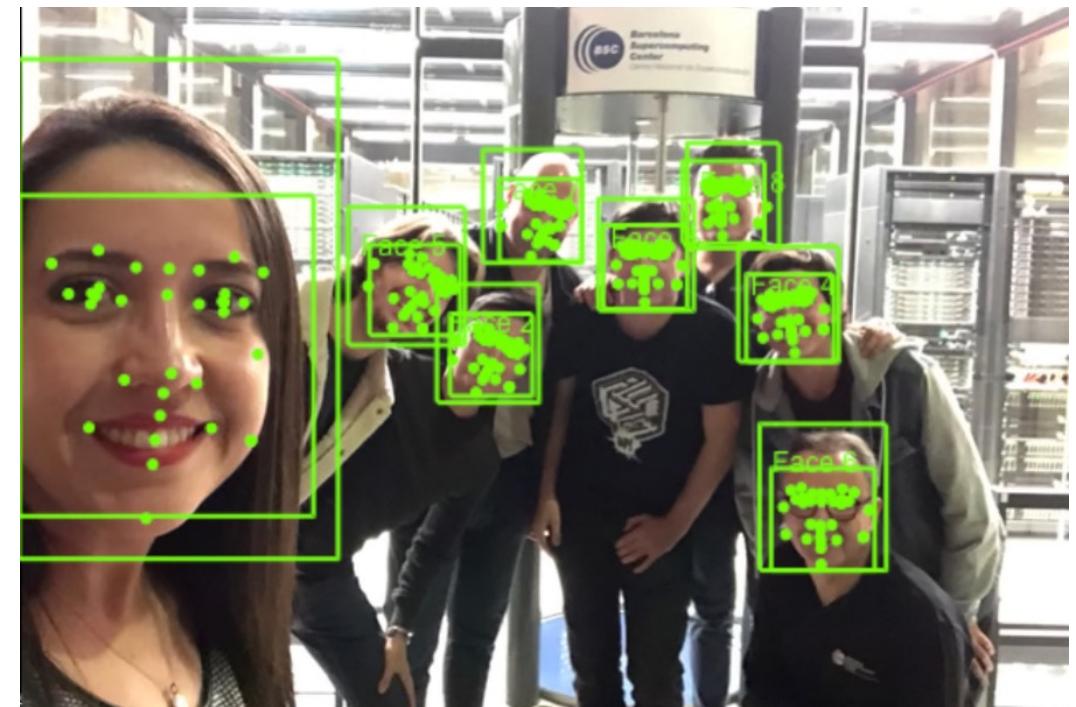
Source: <http://karpathy.github.io/2015/10/25/selfie/>

Jordi TORRES.AI

And also, NOT FACES

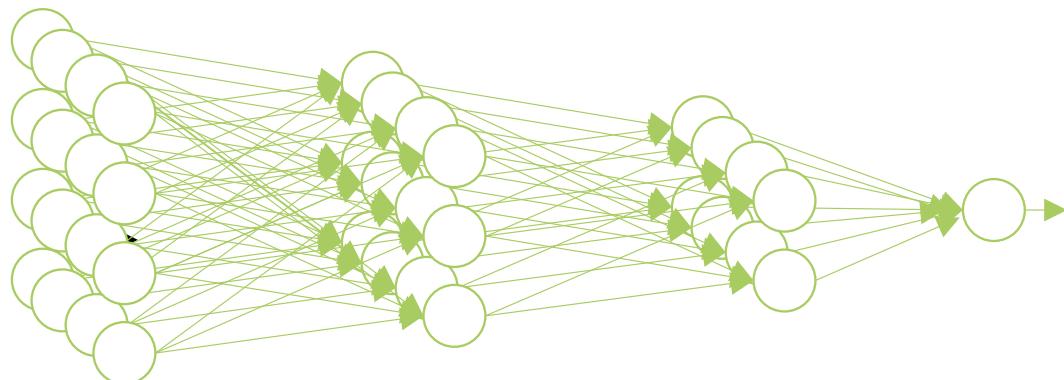


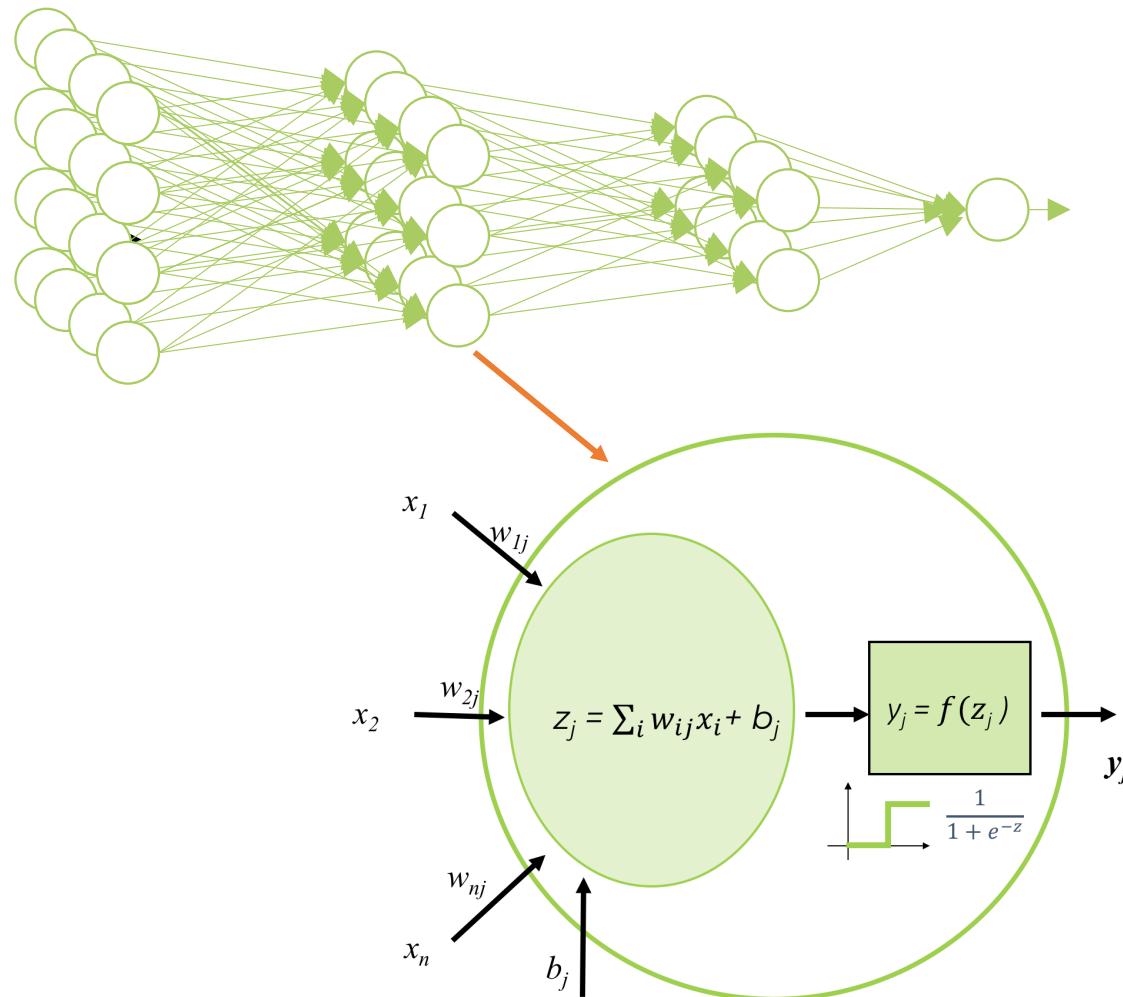
Et voilà!

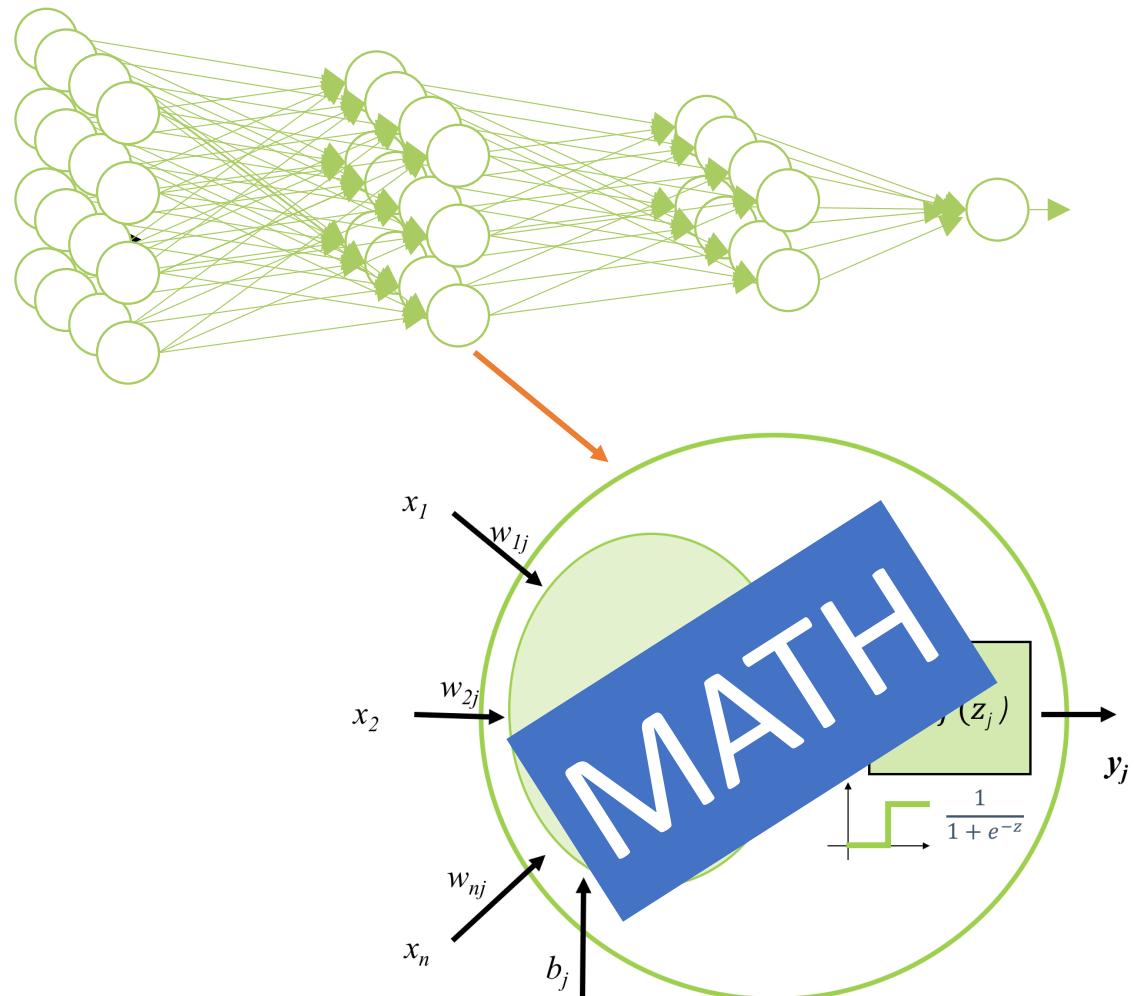


So what's all about?

# Deep Learning: Matematic models





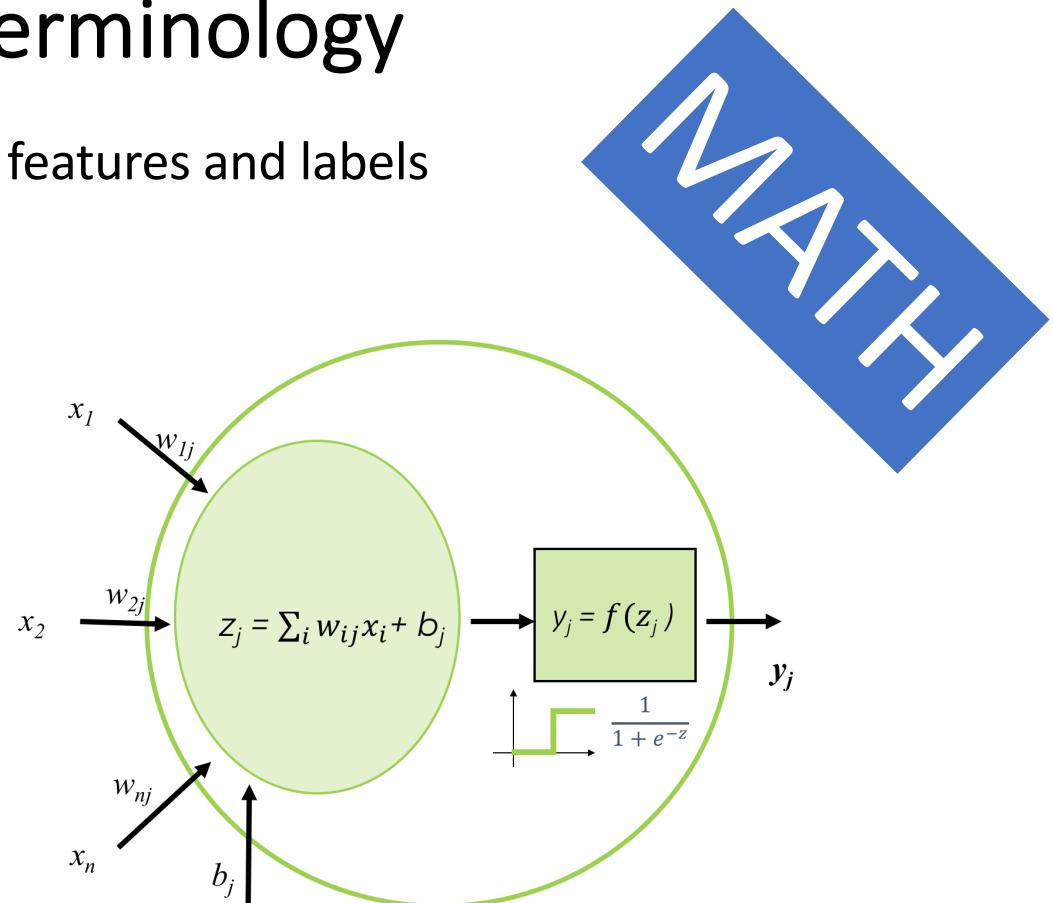


# Basic machine learning terminology

- Model: defines the relation between features and labels

$$y = w\mathbf{x} + b$$

- $y$ : Labels
- $\mathbf{x}$ : Features
- $w$  :weights
- $b$  : bias

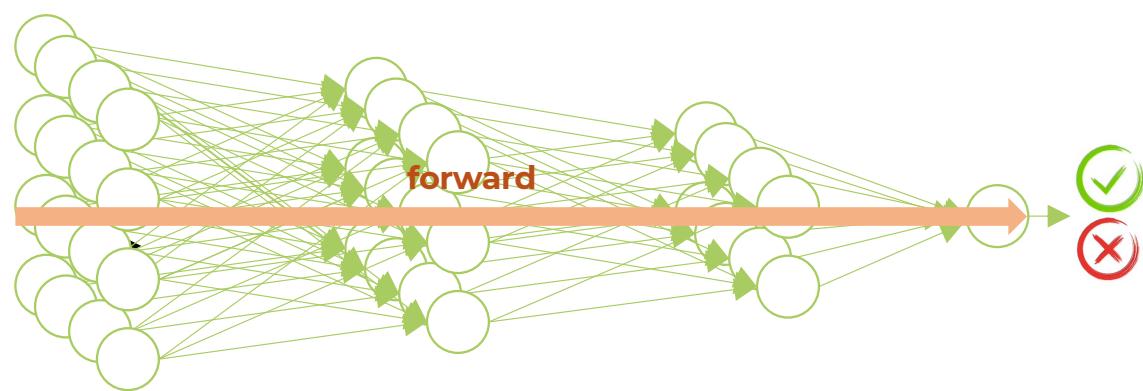


# Basic machine learning terminology

$$y = \sum_i w_i x_i + b$$

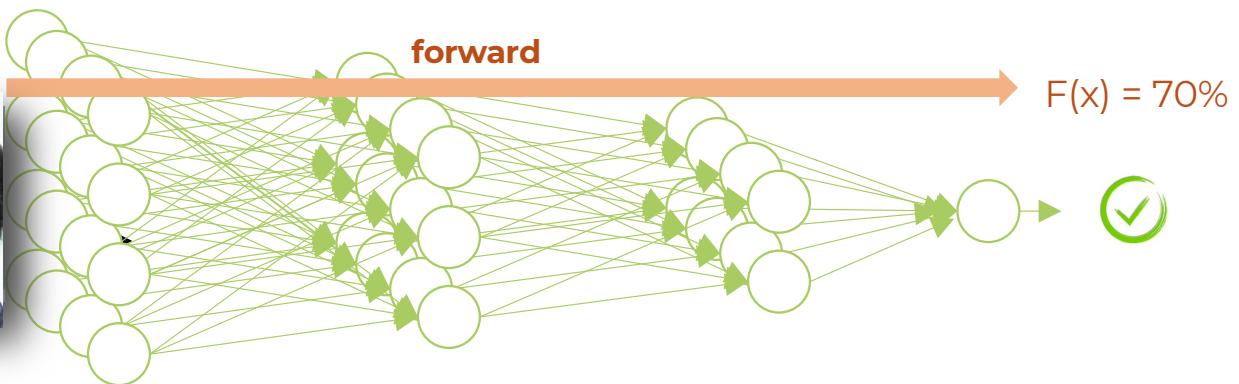
$$y = w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b = \mathbf{w}^T \cdot \mathbf{x} + b$$


$$y = w_0 x_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n = \mathbf{w}^T \cdot \mathbf{x}$$

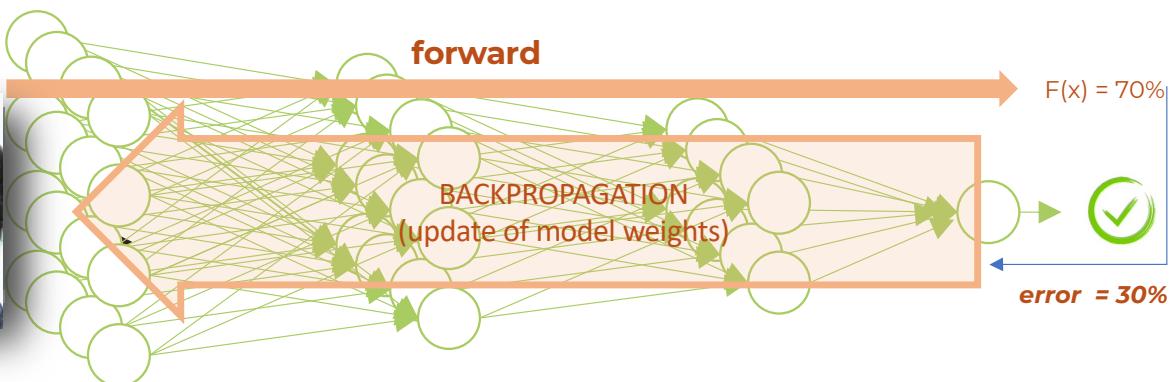


*Deep Learning - Supervised Learning*

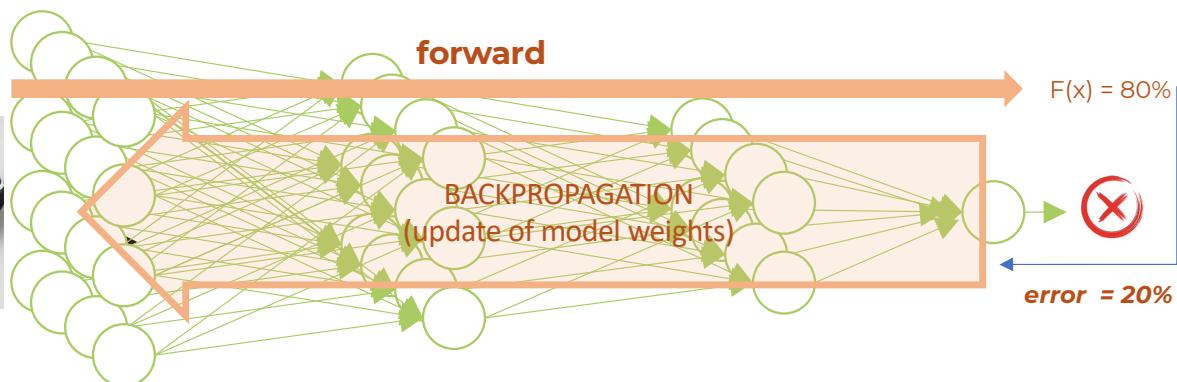
## Fase "TRAINING"



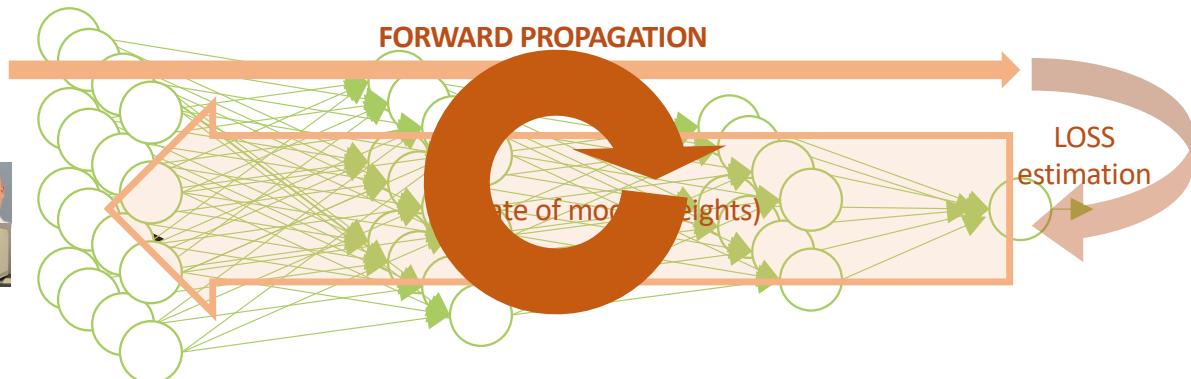
## Fase "TRAINING"



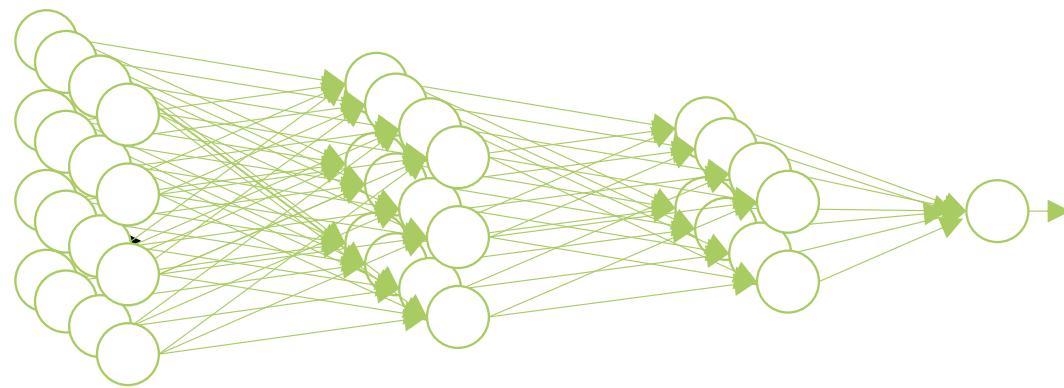
## Fase "TRAINING"



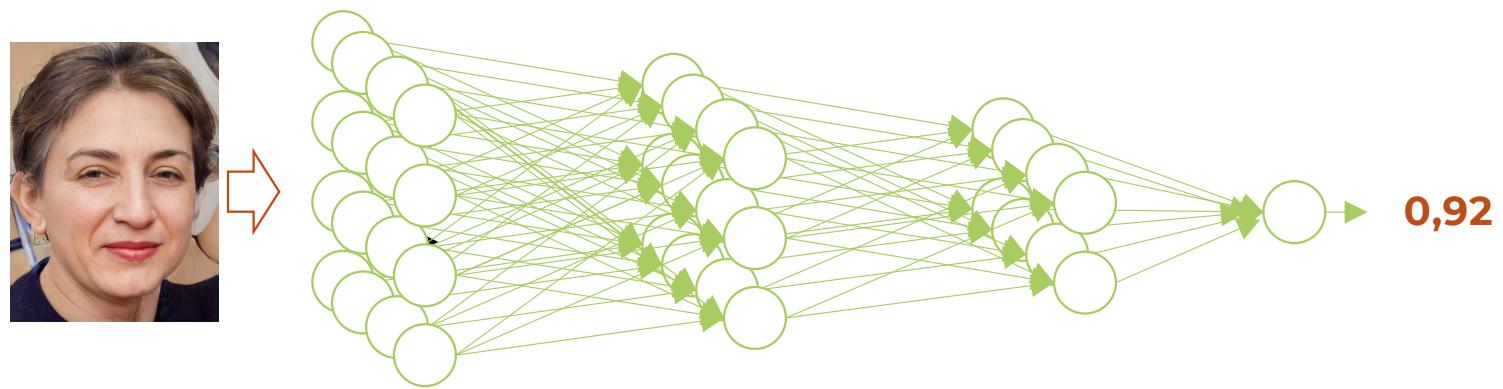
## Fase "TRAINING"



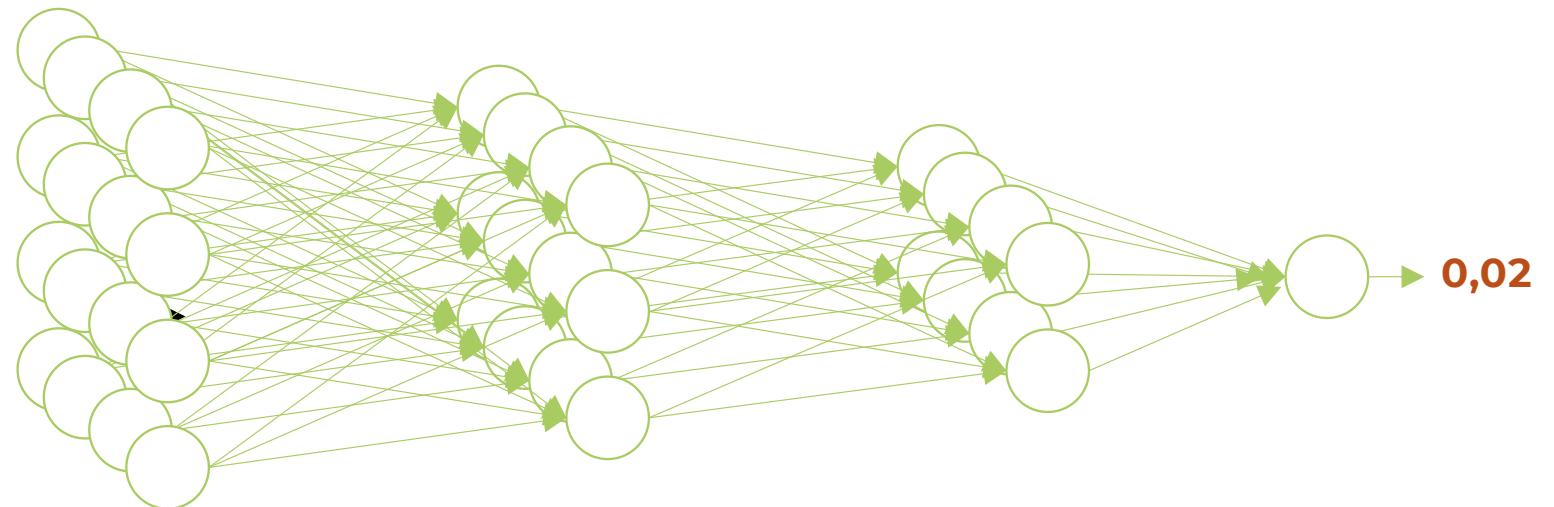
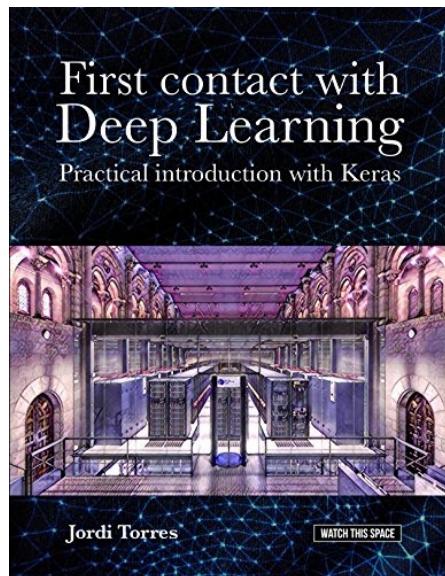
**After "TRAINING" the model is ready for "INFERENCE"**



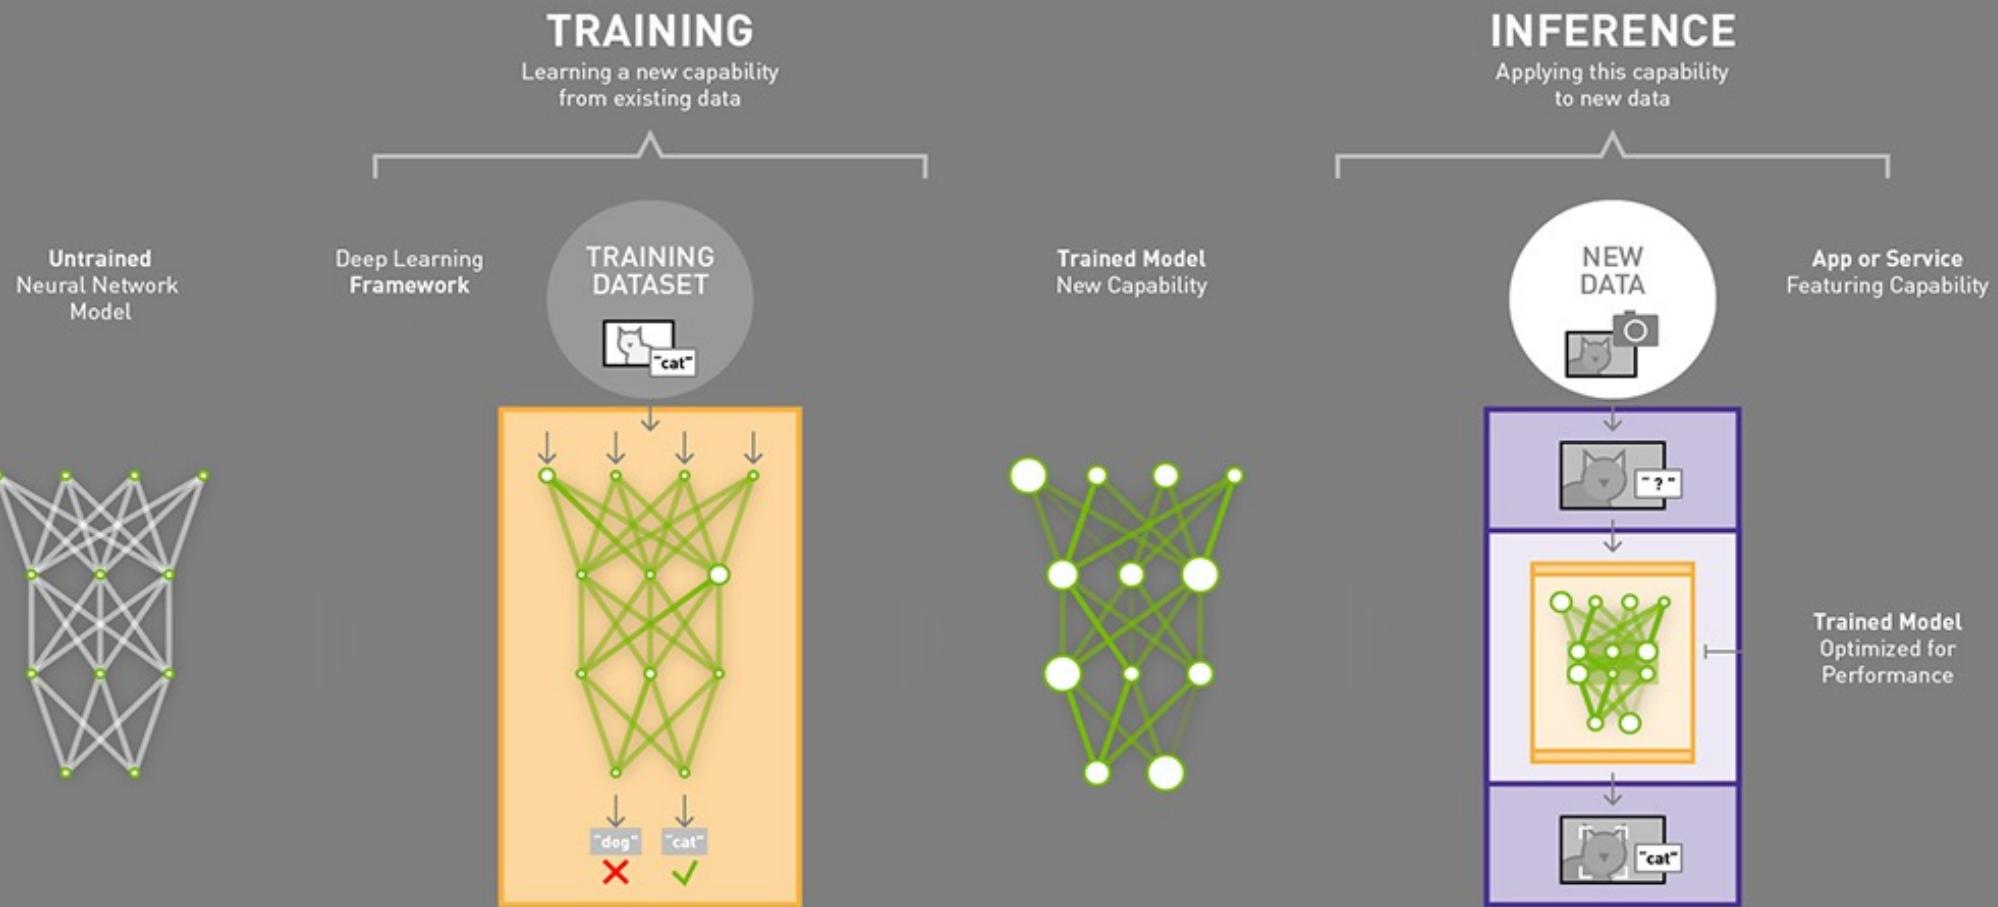
## Fase "INFERENCIA"



## Fase "INFERENCE"



# Deep Learning: Training vs Inference



Source: <https://blogs.nvidia.com/blog/2016/08/22/difference-deep-learning-training-inference-ai/>



Working  
environment?

Jordi **TORRES.AI**



source: [https://www.researchgate.net/publication/349108714\\_DLBench\\_a\\_comprehensive\\_experimental\\_evaluation\\_of\\_deep\\_learning\\_frameworks](https://www.researchgate.net/publication/349108714_DLBench_a_comprehensive_experimental_evaluation_of_deep_learning_frameworks)

# TensorFlow vs PyTorch?

 From your Digest

 **Mike West** · October 5  
Instructor at LogikBot (2020–present) X

**Is PyTorch better than Keras?**

Originally Answered: Which frame is better in the market, Keras or PyTorch?  
I don't think there is a better. However, with the inclusion of Keras into TF 2.0 it's going to be hard for any other frameworks to compete.

I like Keras and have been using it for several years.

You could get an idea of usage in the applied space by searching for PyTorch and then TensorFlow on any job board. I did it on indeed and there are twice as many jobs for TF as there are for PT.

I'd highly recommend Keras on TF.

16.5K views · View Upvoters · View Sharers · Answer requested by Pritika Arjun Kumar

 15    2    3

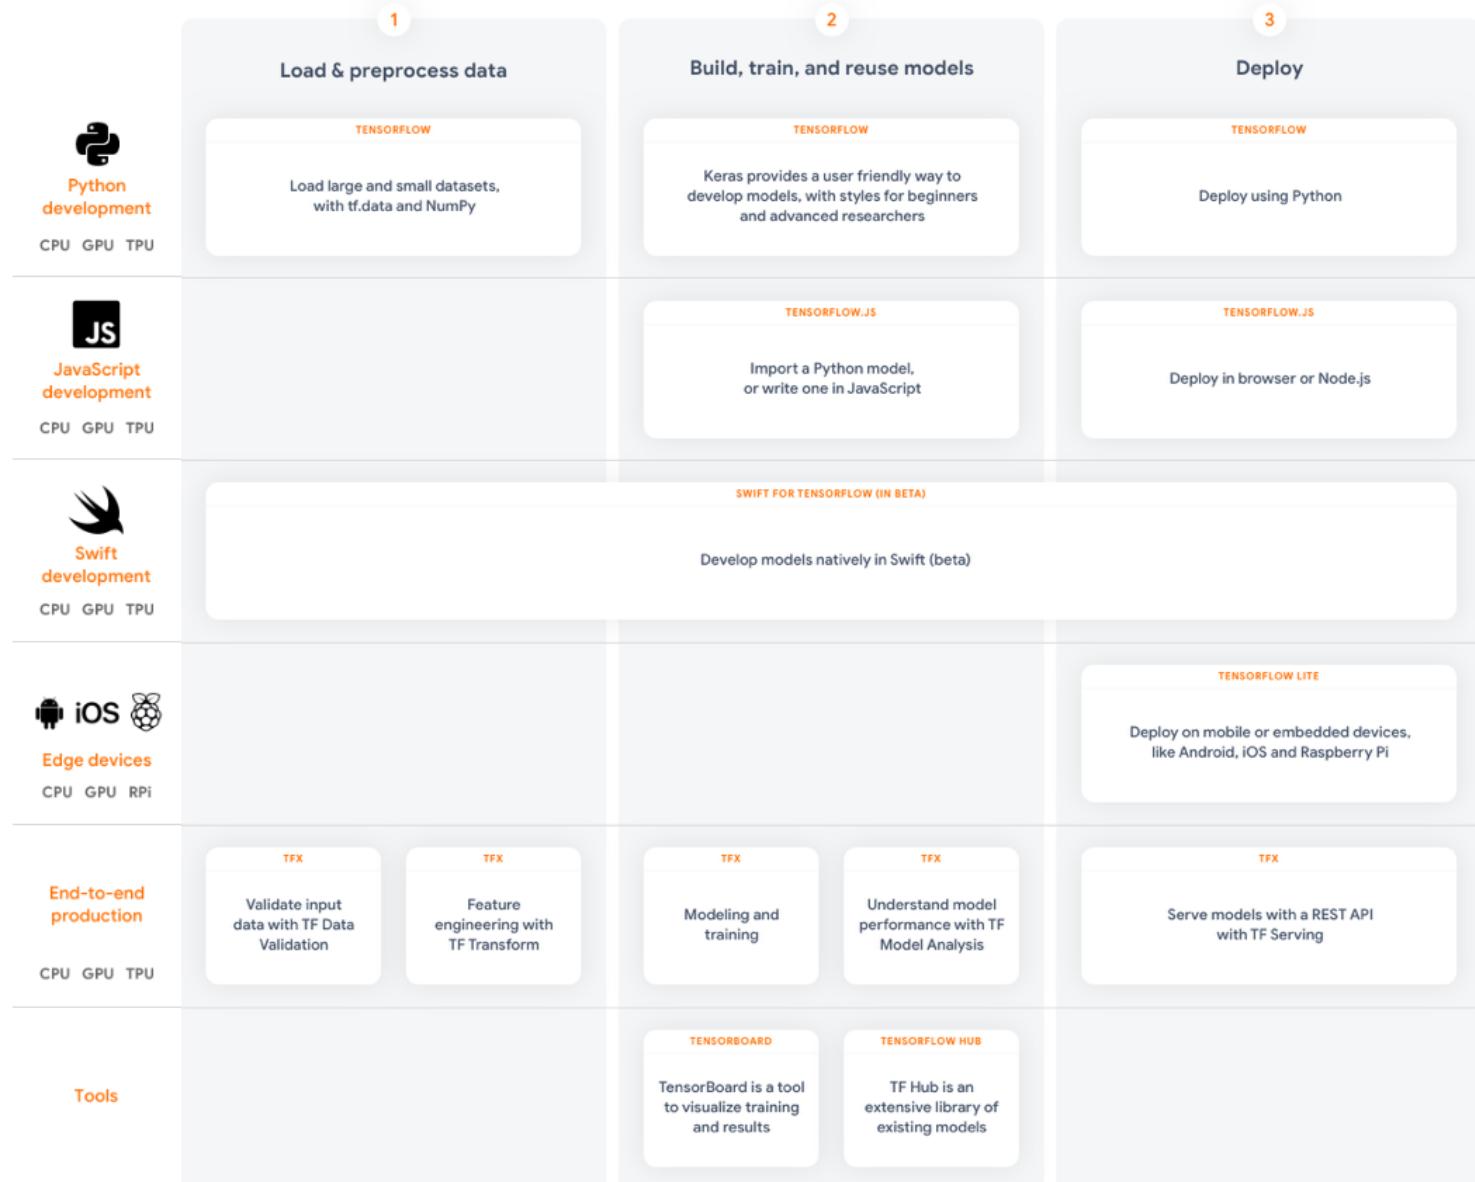
# TensorFlow 2.\*

- During **2019 TensorFlow Dev Summit**, Google announced a major upgrade on the framework → the TensorFlow 2.0 Alpha version.

Latest stable release: 2.20 (13 August 2025)

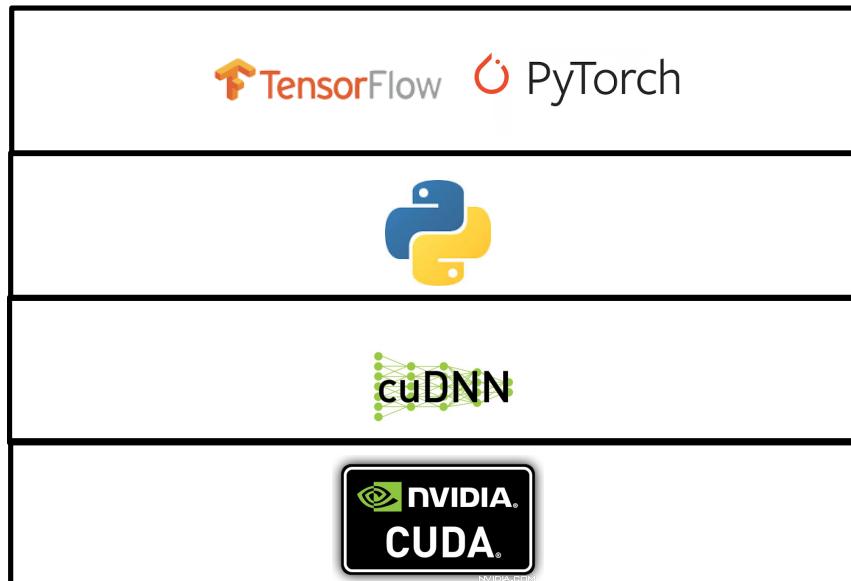
- **TensorFlow 2.0 focuses on simplicity and ease of use:**
  - with updates like eager execution
  - flexible model building on any platform
  - **intuitive higher-level APIs using**  Keras (developed by François Chollet)

# TensorFlow Ecosystem



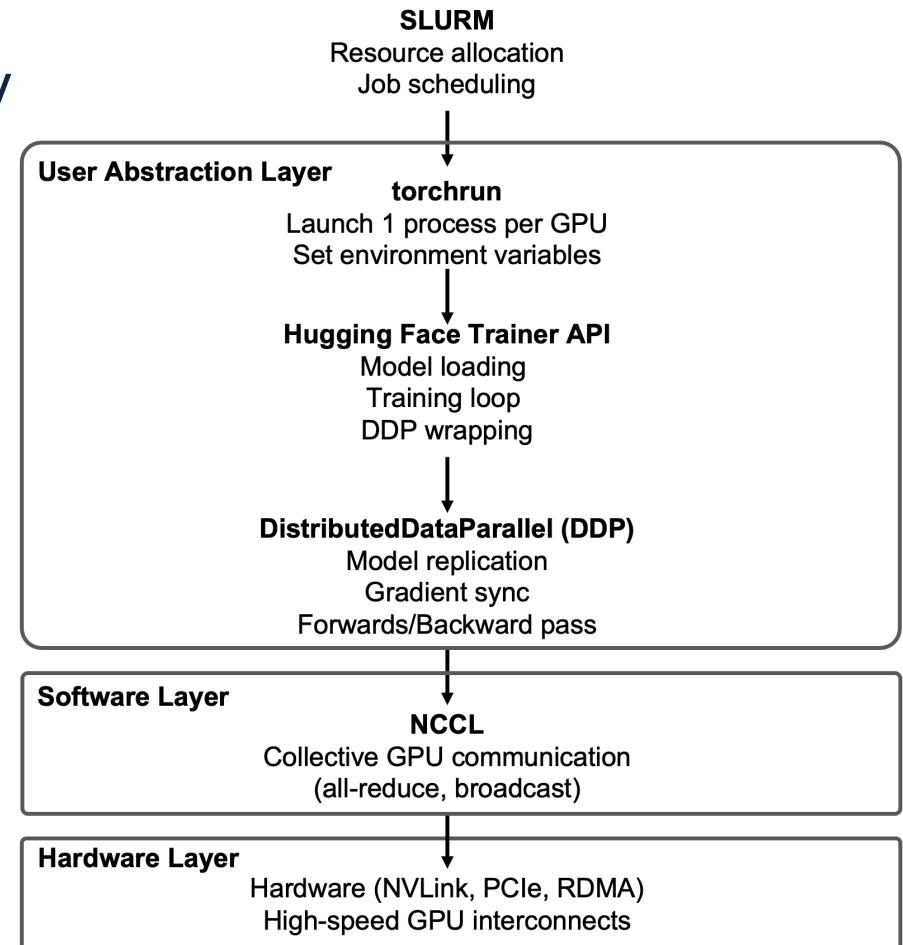
# Deep Learning Software Stack

- Our case study: basic



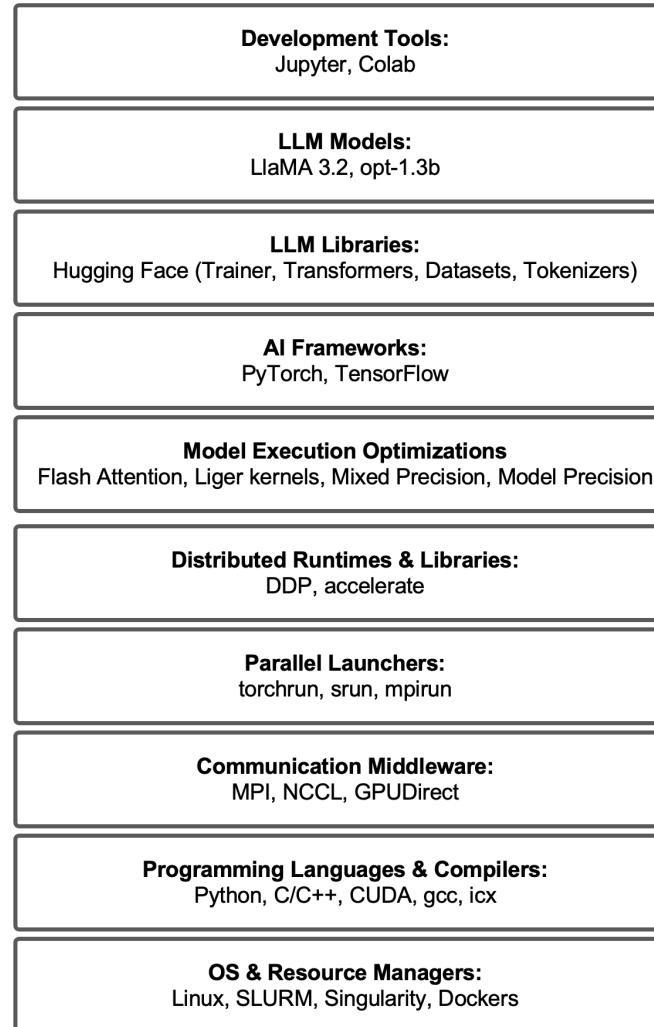
# Deep Learning Software Stack

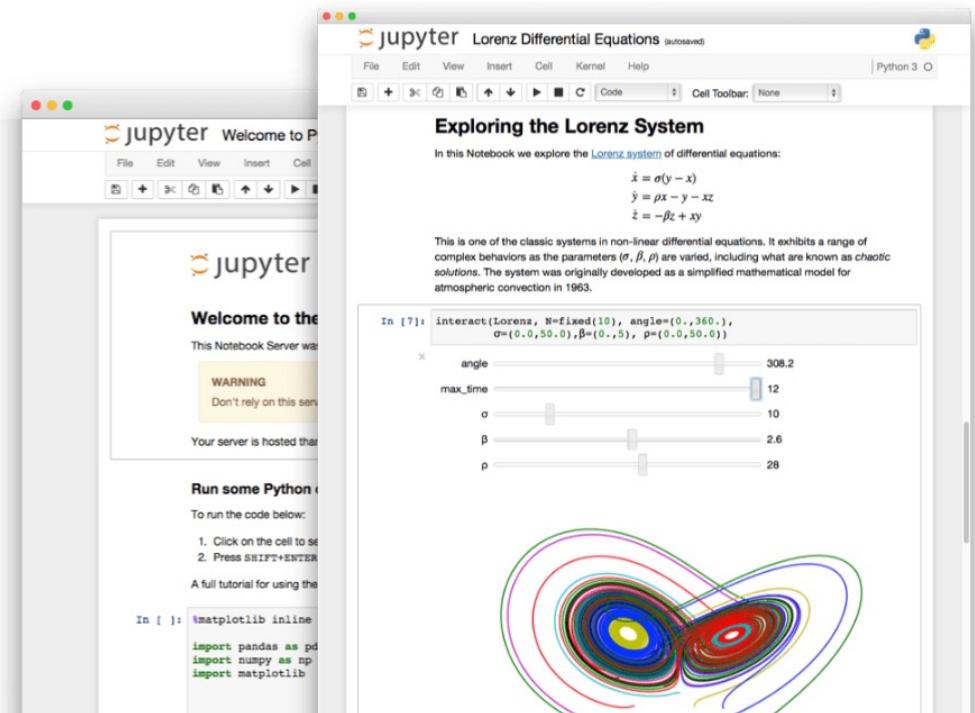
## ■ Our case study: workflow



# Deep Learning Software Stack

## ■ Our case study





## The Jupyter Notebook

The Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live code, equations, visualizations and narrative text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, data visualization, machine learning, and much more.

# Working Environment Colab

## Google Colab

<https://colab.research.google.com/> ▾ Traducir esta página

Colaboratory notebooks are stored in **Google Drive** and can be shared just as you ... **Colab** also supports connecting to a Jupyter runtime on your local machine.

Welcome To Colaboratory - Colab

New Tab

https://colab.research.google.com/notebooks/welcome.ipynb#scrollTo=xitplqMNk\_Hc

... 🌐 ⚡

SHARE Sign in

Welcome To Colaboratory

File Edit View Insert Runtime Tools Help

CODE TEXT

Table of contents Code snippets

Introducing Colaboratory

Getting Started

More Resources

Machine Learning Examples: Seedb

+ SECTION

Run all ⌘/Ctrl+F9

Run before ⌘/Ctrl+F8

Run the focused cell ⌘/Ctrl+Enter

Run selection ⌘/Ctrl+Shift+Enter

Run after ⌘/Ctrl+F10

Interrupt execution ⌘/Ctrl+M I

Restart runtime... ⌘/Ctrl+M R

Restart and run all... ⌘/Ctrl+M A

Reset all runtimes...

Change runtime type

Manage sessions

CONNECT EDITING

Welcome to Colaboratory!

Colaboratory is a Jupyter notebook environment that requires no setup and runs entirely in the cloud.

You can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser.

Watch video

Intro to Google Colab

Watch later Share

Coding TensorFlow

The screenshot shows the Google Colaboratory interface. On the left, there's a sidebar with links like 'Introducing Colaboratory', 'Getting Started', 'More Resources', and 'Machine Learning Examples: Seedb'. The main area has tabs for 'CODE' and 'TEXT'. A 'Table of contents' section is visible. The top navigation bar includes 'File', 'Edit', 'View', 'Insert', 'Runtime', 'Tools', and 'Help'. A context menu is open over the 'Runtime' tab, listing options like 'Run all', 'Run before', etc., with 'Change runtime type' highlighted by a blue oval. The main content area displays a video thumbnail for 'Intro to Google Colab' with a play button and the text 'Coding TensorFlow'. The URL in the address bar is https://colab.research.google.com/notebooks/welcome.ipynb#scrollTo=xitplqMNk\_Hc.

Welcome To Colaboratory - Colab

New Tab

https://colab.research.google.com/notebooks/welcome.ipynb#scrollTo=xitplqMNk\_Hc

CO Welcome To Colaboratory

File Edit View Insert Runtime Tools Help

CODE TEXT CELL CELL COPY TO DRIVE

CONNECT EDITING

Table of contents Code snippets Files

Introducing Colaboratory

Getting Started

More Resources

Machine Learning Examples: Seedbank

SECTION

# Welcome to Colaboratory!

Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud.

With Colaboratory you can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser.

### Notebook settings

Runtime type: Python 3

Hardware accelerator: GPU

Omit code cell output when saving this notebook

CANCEL SAVE

Introducing Colaboratory

This 3-minute video shows you how to get started with Colaboratory.

Get started

Introducing Colaboratory

Coding TensorFlow

Welcome To Colaboratory - Colab

New Tab

https://colab.research.google.com/notebooks/welcome.ipynb?pli=1#scrollTo=xitplqMNk\_Hc

File Edit View Insert Runtime Tools Help

New Python 3 notebook

New Python 2 notebook

Open notebook... %/Ctrl+O

Upload notebook...

Rename...

Move to trash

Save a copy in Drive...

Save a copy as a GitHub Gist...

Save a copy in GitHub...

Save %/Ctrl+S

Save and pin revision %/Ctrl+M S

Revision history

Download .ipynb

Download .py

Update Drive preview

Print %/Ctrl+P

Y TO DRIVE

RAM Disk EDITING

SHARE L

# Welcome to Colaboratory!

Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud.

In Colaboratory you can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser.

## Introducing Colaboratory

This 3-minute video gives an overview of the key features of Colaboratory:

Get started with Google Colaboratory (Coding...)

Watch later Share

## Intro to Google Colab

Coding TensorFlow



# Steps to create your first neural network

Jordi **TORRES.AI**

# Steps to create a model with TensorFlow

1

Load Data &  
Prepare Data

2

Define  
Model

3

Compile  
Model

4

Fit Model

5

Evaluate  
Model

6

Use the  
Model

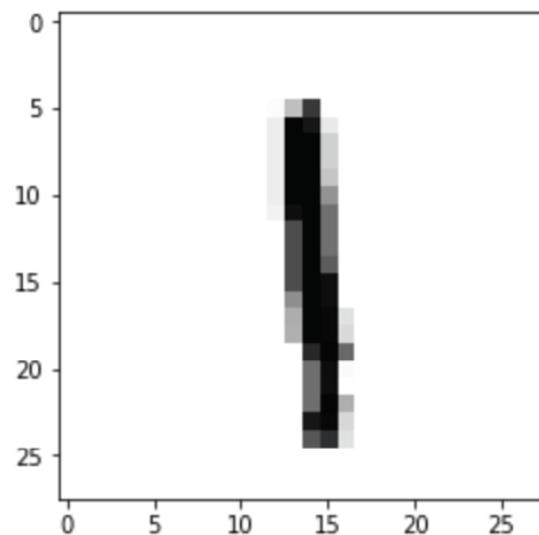
# 1. Load and prepare data

- Case study: “The MNIST database”
  - Dataset of handwritten digits classification
  - 60,000 28x28 grayscale images of the 10 digits, along with a test set of 10,000 images.



# 1. Load and prepare data

- Features (Images): matrix of 28x28 pixels with values [0, 255]
- Labels: values [0, 9]



# 1. Load and prepare data: Example

- Feature (Image): matrix of 28x28 pixels with values [0, 255]

- Label: 1

# 1. Load and prepare data

```
In [2]: mnist = tf.keras.datasets.mnist  
(x_train, y_train), (x_test, y_test) = mnist.load_data(path='/gpfs/projects/nct00/nct001')
```

## 2. Define model

```
In [24]: from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(10, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

### 3. Compile model

```
In [26]: model.compile(loss="categorical_crossentropy",
                      optimizer="sgd",
                      metrics = ['accuracy'])
```

## 4. Fit model

```
[24] model.fit(x_train, y_train, epochs=5)

→ Epoch 1/5
60000/60000 [=====] - 6s 98us/sample - loss: 0.4039 - accuracy: 0.8930
Epoch 2/5
60000/60000 [=====] - 6s 94us/sample - loss: 0.1958 - accuracy: 0.9441
Epoch 3/5
60000/60000 [=====] - 5s 88us/sample - loss: 0.1421 - accuracy: 0.9590
Epoch 4/5
60000/60000 [=====] - 5s 86us/sample - loss: 0.1101 - accuracy: 0.9688
Epoch 5/5
60000/60000 [=====] - 5s 86us/sample - loss: 0.0878 - accuracy: 0.9750
<tensorflow.python.keras.callbacks.History at 0x7fbdb8af52ba8>
```

## 5. Evaluate model

```
[25] test_loss, test_acc = model.evaluate(x_test, y_test)
    print ('Test accuracy:', test_acc)

→ 10000/10000 [=====] - 0s 48us/sample - loss: 0.0955 - accuracy: 0.9716
Test accuracy: 0.9716
```

## 6. Use the model

```
[31] predictions = model.predict(x_test)
     print (predictions)

[ [1.26404475e-05 2.49308573e-06 4.45205660e-04 ... 9.93939281e-01
  2.26089851e-05 3.87840439e-04]
  [1.40489414e-04 5.95517922e-03 9.92044926e-01 ... 7.36168673e-08
  8.24517483e-05 5.90428044e-08]
  [4.98437309e-08 9.97506320e-01 1.01395021e-03 ... 9.72060545e-04
  2.70858494e-04 1.64484100e-05]
  ...
  [1.58638017e-07 4.75961940e-07 7.93745471e-07 ... 1.92583975e-04
  1.82772157e-04 4.21503466e-03]
  [2.74540548e-06 5.57959866e-05 4.68355069e-07 ... 3.29460909e-06
  5.54329483e-03 1.57847660e-06]
  [2.11494735e-05 1.39356416e-06 7.07783474e-05 ... 2.89122028e-07
  5.49738729e-07 9.28026580e-08]]
```

# Live Demo: Get started with Deep Learning basics

1. Open Colab
2. Upload

<https://github.com/jorditorresBCN/HPC4AIbook/blob/main/Chapter.07/PracticalIntroductionToDeepLearningBasics.ipynb>

# To get started, import the TensorFlow library into your program:

```
[1]
import tensorflow as tf
from tensorflow import keras

import numpy as np
import matplotlib.pyplot as plt

print(tf.__version__)

2.19.0
```

# Steps to create a model with TensorFlow

1

Load Data &  
Prepare Data

2

Define  
Model

3

Compile  
Model

4

Fit Model

5

Evaluate  
Model

6

Use the  
Model

# 1. Load and prepare data

- Case study: “The MNIST database”
  - Dataset of handwritten digits classification
  - 60,000 28x28 grayscale images of the 10 digits, along with a test set of 10,000 images.

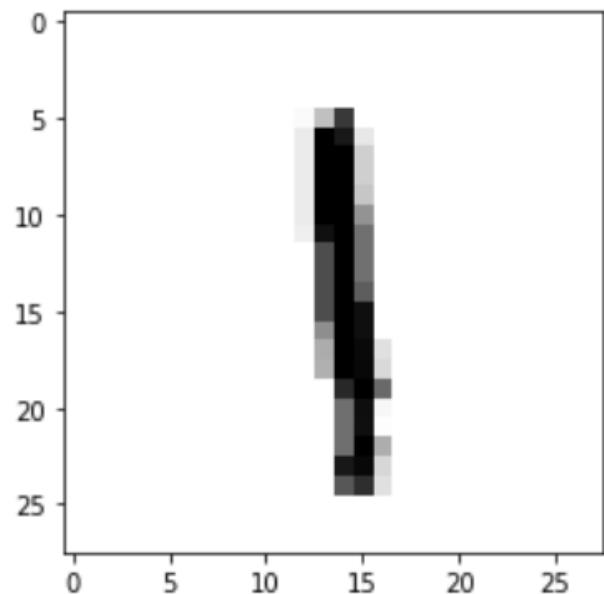


```
mnist = tf.keras.datasets.mnist  
(x_train, y_train), (x_test, y_test) = mnist.load_data(path='/gpfs/projects/nct00/nct00')
```

MN5: path='/gpfs/.../mnist.npz'

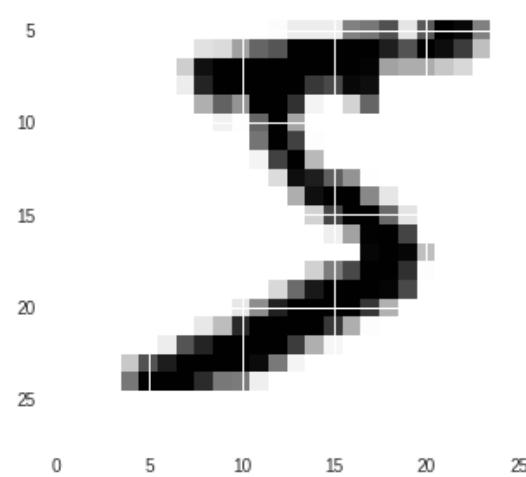
```
In [3]: import matplotlib.pyplot as plt  
plt.imshow(x_train[8], cmap=plt.cm.binary)
```

```
Out[3]: <matplotlib.image.AxesImage at 0x7fff30123e10>
```



```
[9] import matplotlib.pyplot as plt  
plt.imshow(x_train[0], cmap=plt.cm.binary)
```

↳ <matplotlib.image.AxesImage at 0x7fb96975b70>

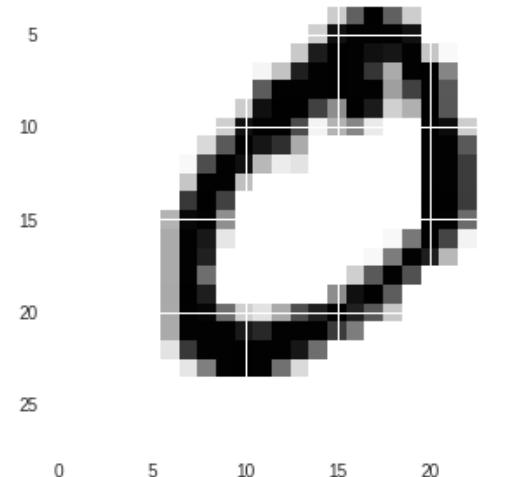


```
[11] print(y_train[0])
```

↳ 5

```
[12] import matplotlib.pyplot as plt  
plt.imshow(x_train[1], cmap=plt.cm.binary)
```

↳ <matplotlib.image.AxesImage at 0x7fb96951630>



```
▶ print(y_train[1])
```

↳ 0

# 1. Load and prepare data

- Normalizing the data

- Original images of  $28 \times 28$  pixels whose values range from [0, 255] of type uint8.

```
In [4]: print(y_train[8])
```

```
1
```

```
In [5]: print(x_train.ndim)
```

```
3
```

```
In [6]: print(x_train.shape)
```

```
(60000, 28, 28)
```

```
In [7]: print(x_train.dtype)
```

```
uint8
```

# 1. Load and prepare data

- Normalizing the data
  - But it is usual to scale : in the example to [0, 1].

```
[8] x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')

    x_train /= 255
    x_test /= 255
```

# 1. Load and prepare data

- reshape 2D → 1D

```
[9] print(x_train.shape)  
print(x_test.shape)
```

```
(60000, 28, 28)  
(10000, 28, 28)
```

```
[10] x_train = x_train.reshape(60000, 784)  
x_test = x_test.reshape(10000, 784)
```

```
[11] print(x_train.shape)  
print(x_test.shape)
```

```
(60000, 784)  
(10000, 784)
```

# 1. Load and prepare data

- categorize data (one-shot):

```
[13] print(y_train[0])
```

```
5
```

```
[14] print(y_train.shape)
```

```
(60000,)
```

```
[16] from tensorflow.keras.utils import to_categorical
```

# 1. Load and prepare data

- categorize data (one-shot):

```
[17] y_train = to_categorical(y_train, num_classes=10)
     y_test = to_categorical(y_test, num_classes=10)
```

```
[18] print(y_test[0])
     print(y_train[0])
```

```
[0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]
[0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
```

```
[19] print(y_train.shape)
     print(y_test.shape)
```

```
(60000, 10)
(10000, 10)
```

## 2. Define model

```
In [24]: from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(10, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

## Model: core data structure

- [keras.models.Sequential](#) class is a wrapper for the neural network model
- Models in Keras are defined as a sequence of layers.

## 2. Define model

```
In [25]: model.summary()
```

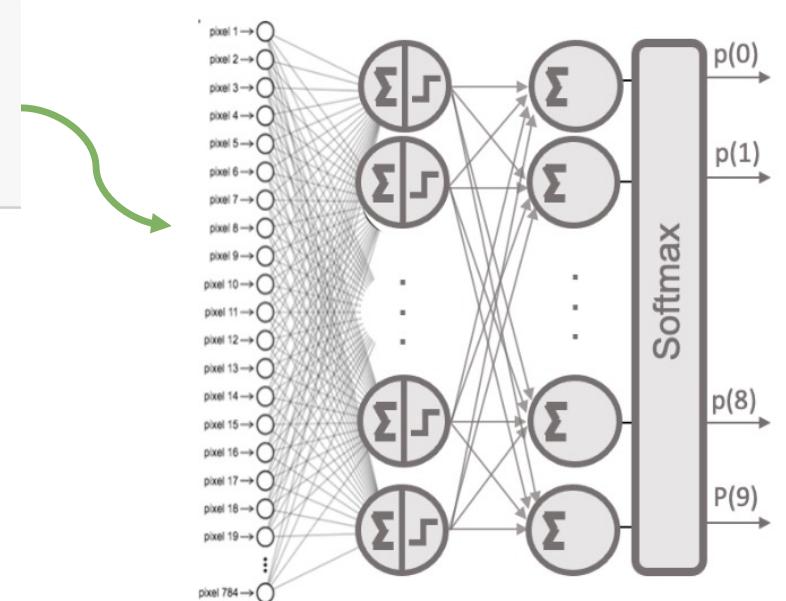
Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	7850
dense_1 (Dense)	(None, 10)	110
<hr/>		
Total params: 7,960		
Trainable params: 7,960		
Non-trainable params: 0		
<hr/>		

- **Keras will automatically infer the shape of all layers after the first layer.**
- This means you only have to set the input dimensions for the first layer.

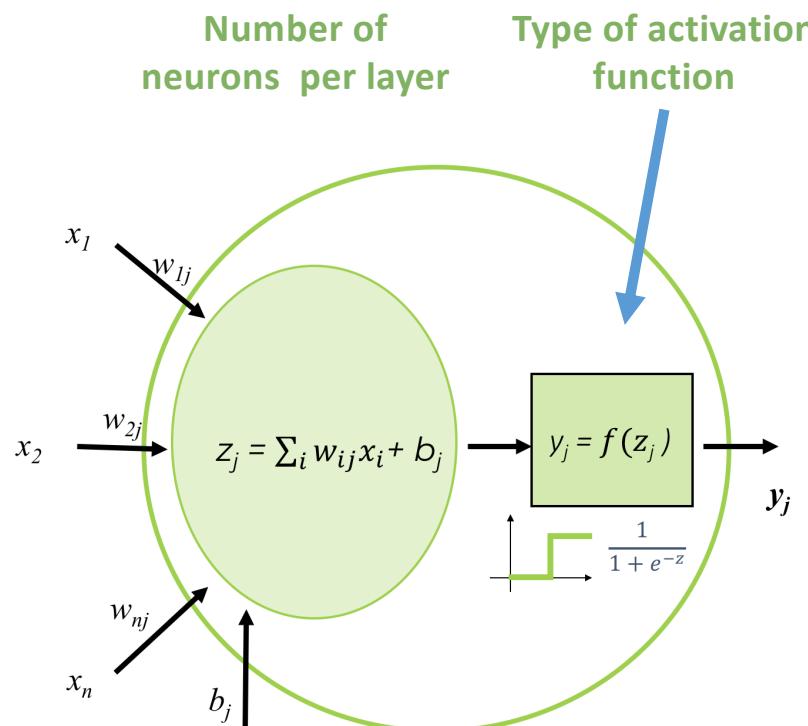
## 2. Define model

```
In [24]: from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(10, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```



## 2. Define model

```
In [24]: from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(10, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```



## 2. Define model

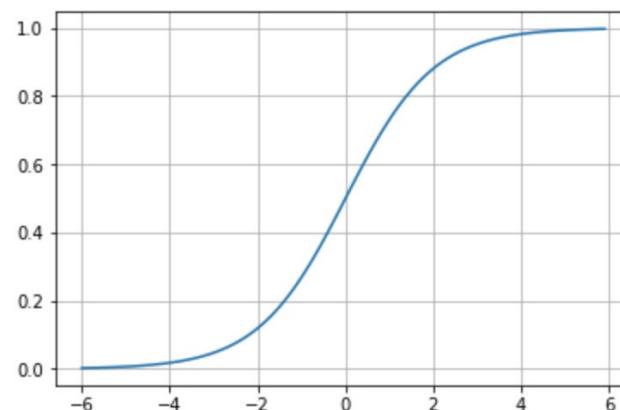
```
In [24]: from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(10, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

$$y = \frac{1}{1+e^{-z}}$$

The sigmoid function allows us to reduce extreme or atypical values without eliminating them.

converts independent variables of almost infinite range into simple probabilities between 0 and 1.

Most of the output will be very close to the extremes of 0 or 1, as we have already commented.



## 2. Define model

```
In [24]: from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(10, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```



Softmax?

## 2. Define model

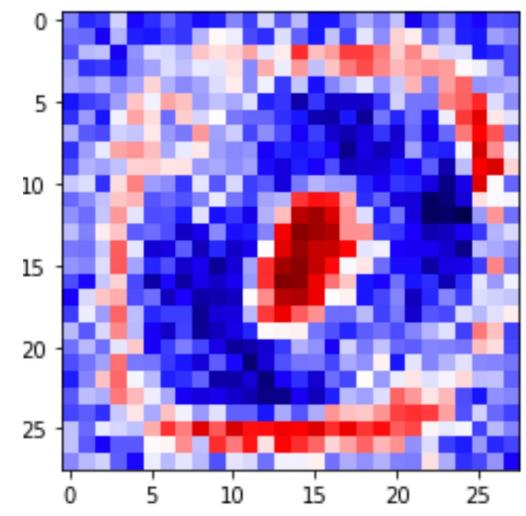
- The *softmax* function has two main steps:
  - first, the “**evidences**” for an image belonging to a certain label are computed,
  - and later the evidences are converted into **probabilities** for each possible label.
- Example: MNIST database



## Evidence of belonging

- An approach to measure the evidence that a certain image belongs to a particular class is to make a weighted sum of the evidence of belonging to each of its pixels to that class.

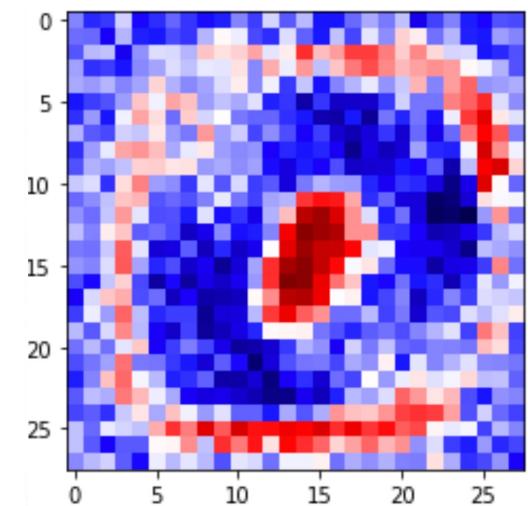
To explain the idea I will use a visual example ->



## Evidence of belonging

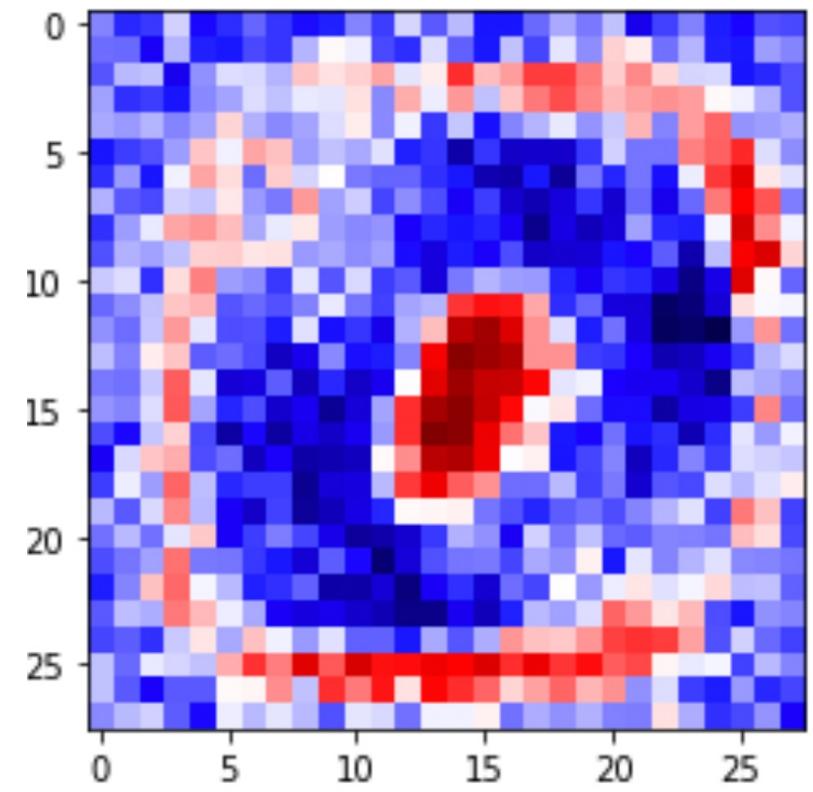
Let's suppose that we already have the model learned for the number zero (28x28):

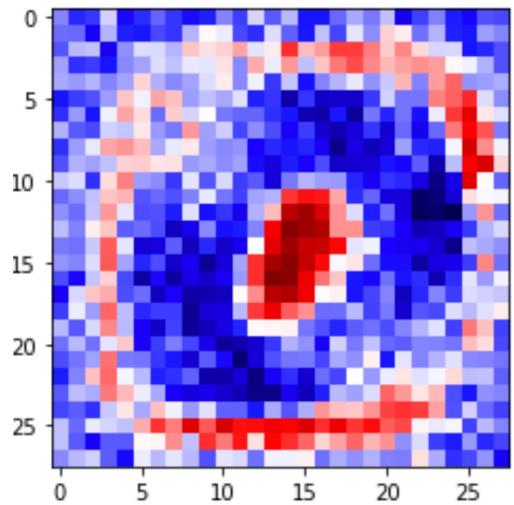
- Pixels in **red** represent negative weights (i.e., reduce the evidence that it belongs),
- Pixels in **blue** represent positive weights (the evidence of which is greater increases).
- The **black** color represents the neutral value.



# Evidence of belonging

- Let's imagine that we **trace a zero** over it.
  - **In general, the trace of our zero would fall on the blue zone**
  - It is quite evident that if our stroke goes over the red zone, it is most likely that we are not writing a zero;
  - therefore, using a metric based on **adding if we pass through** the **blue** zone and **subtracting if we pass through** the **red** zone seems reasonable.

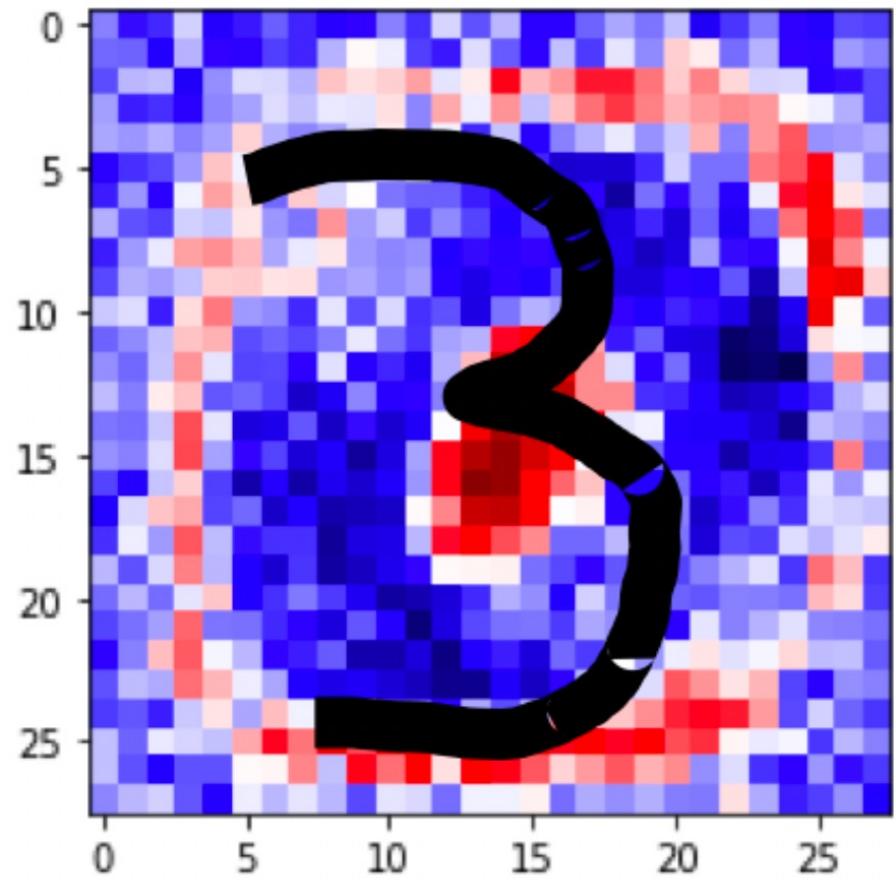




```
[ [ 22. -21. 11. -8. 0. -17. -4. -19. 12. 21. -10. 12. -20. -19. 11. -21. 14. -13. -21. -10. 17. -5. -11. 15. -16. -19. -21. 4.]  
[ 10. 1. 2. 5. 20. -5. 16. -19. -1. -16. -9. 25. -9. -11. -48. 7. -10. -14. -15. -11. 6. -11. 17. -21. 20. -17. -20. -18.]  
[ 1. 17. 16. 3. -19. -16. 3. -15. -33. -46. -56. -53. -42. -56. -99. -88. -104. -98. -86. -79. -42. -54. -51. -37. -4. 13. 21. 17.]  
[ -4. -7. 5. 5. -27. -9. -17. -22. -12. -29. -46. -33. -50. -41. -31. -50. -13. -44. -82. -49. -66. -55. -78. -52. -34. -4. -19. -14.]  
[ 11. -4. -7. 13. -60. -44. -68. -50. -38. -30. -23. -23. -1. -10. -21. 13. -24. -13. -6. -19. -21. -15. -21. -58. -61. -34. -63. -33.]  
[ 16. -8. -30. -21. -72. -46. -45. -38. -28. -22. -23. -21. -13. 3. 43. 32. 15. 19. 21. -1. -16. 7. -21. -44. -81. -92. -66. -14.]  
[ -21. 10. -31. -49. -46. -33. -36. -36. -10. -36. -15. 4. -6. 9. 32. 26. 24. 55. 44. 16. 29. 6. 33. -2. -82. -142. -58. -24.]  
[ -7. -26. -23. -46. -89. -16. -3. -37. -37. -31. -1. 18. 17. 24. 9. 16. 28. 42. 36. 38. 22. 16. 16. 19. -50. -107. -89. -12.]  
[ 29. -38. -21. -23. -62. -50. -33. -27. -33. -6. -14. -2. 7. -2. 10. 28. 42. 62. 60. 23. 1. 18. 32. 26. -43. -155. -98. -39.]  
[ -8. -13. -1. -42. -50. -24. -43. -41. -24. 8. 8. 1. 19. 34. 18. 7. 21. 42. 47. 43. 19. 7. 34. 46. 6. -145. -89. -33.]  
[ 2. -10. 9. -27. -32. -16. -10. -11. -10. -9. -24. -15. 16. 7. 2. -28. -24. 27. 44. 28. 11. 42. 22. 40. 23. -96. -94. -34.]  
[ -9. -1. 1. -53. -62. -39. -14. -18. -11. -23. -31. -1. -15. -22. -71. -92. -100. -47. -7. 9. 8. 19. 54. 40. 60. -71. -95. -49.]  
[ 10. 8. -32. -22. -42. -13. 21. 4. -7. 5. 17. 13. -14. -65. -112. -156. -118. -64. -31. 8. 23. 25. 57. 68. 81. -17. -39. -27.]  
[ 4. 1. 5. -22. -14. -12. 27. 11. -3. 15. 11. 15. -4. -68. -132. -165. -133. -56. -9. -20. -17. 15. 46. 47. 62. -12. -61. -38.]  
[ 12. -24. -12. -38. -38. 29. 24. 57. 27. 12. 21. 5. -36. -103. -152. -139. -112. -70. -13. -22. -7. 12. 49. 67. 63. -15. -40. -6.]  
[ -18. -5. -51. -44. 8. 34. 25. 44. 36. 53. 54. 1. -45. -118. -181. -130. -104. -55. -14. -22. 14. 11. 22. 39. 32. -27. -43. -18.]  
[ -17. -20. -20. -54. 0. 30. 30. 47. 26. 60. 45. -9. -72. -130. -129. -99. -56. -3. 1. 19. -8. 18. 5. 7. 9. -49. -44. -20.]  
[ 20. -19. -33. -68. 0. 13. 25. 45. 13. 47. 41. -7. -74. -144. -119. -89. -48. -7. 2. -6. 20. 6. 26. 2. 9. -39. -36. -10.]  
[ -23. 14. -40. -85. -11. -1. 4. 34. 17. 62. 73. 8. -73. -92. -92. -38. -12. 2. -18. 3. -2. 20. 22. 1. -38. -30. -18. -29.]  
[ 9. 17. -11. -74. -31. 3. 27. 2. 34. 48. 37. 31. 9. -51. -12. -19. -8. 7. -33. -5. -30. -7. 8. -11. -37. -47. -47. -37.]  
[ -22. -4. -32. -54. -19. 12. 21. -8. 23. 16. 39. 55. 46. -3. 11. -16. -0. -28. -19. -14. -1. 14. -18. -10. -57. -30. 5. 1.]  
[ -12. -25. -30. -42. -11. -7. 13. 5. 21. 10. 50. 54. 18. 34. 4. 20. -7. -22. -15. -49. -7. -30. -4. -19. -52. -34. 9. 10.]  
[ -17. 2. -5. -31. -75. -15. 20. -3. 19. 11. 49. 49. 38. 48. 21. 10. -7. -12. -20. -16. -41. -27. -7. -17. -18. -21. -4. 10.]  
[ -3. -13. -26. -56. -77. -46. -24. -1. 12. 36. 49. 47. 44. 51. 61. 24. 8. 4. -36. -51. -52. -52. -27. -75. -49. -15. 1. -9.]  
[ 14. -16. -45. -28. -31. -45. -64. -52. -30. -18. 4. 16. 24. 7. 5. -13. -18. -64. -84. -71. -99. -99. -66. -56. -42. -5. 4. 9.]  
[ -22. -14. 3. -15. -29. -46. -51. -75. -103. -105. -129. -119. -125. -121. -142. -119. -134. -117. -102. -86. -76. -80. -30. -65. -34. -0. -20. -7.]  
[ -13. 6. 3. -11. 3. -52. -21. -32. -73. -100. -87. -98. -89. -85. -98. -72. -93. -69. -82. -39. -53. -62. -6. 4. -22. 0. 18. 10.]  
[ -9. 15. 3. 15. -8. -33. 12. -24. -25. -14. -40. -24. -46. -45. -38. -51. -13. -19. -41. -10. -45. -25. -14. -6. 12. 21. -13. -19. ]]
```

# Evidence of belonging

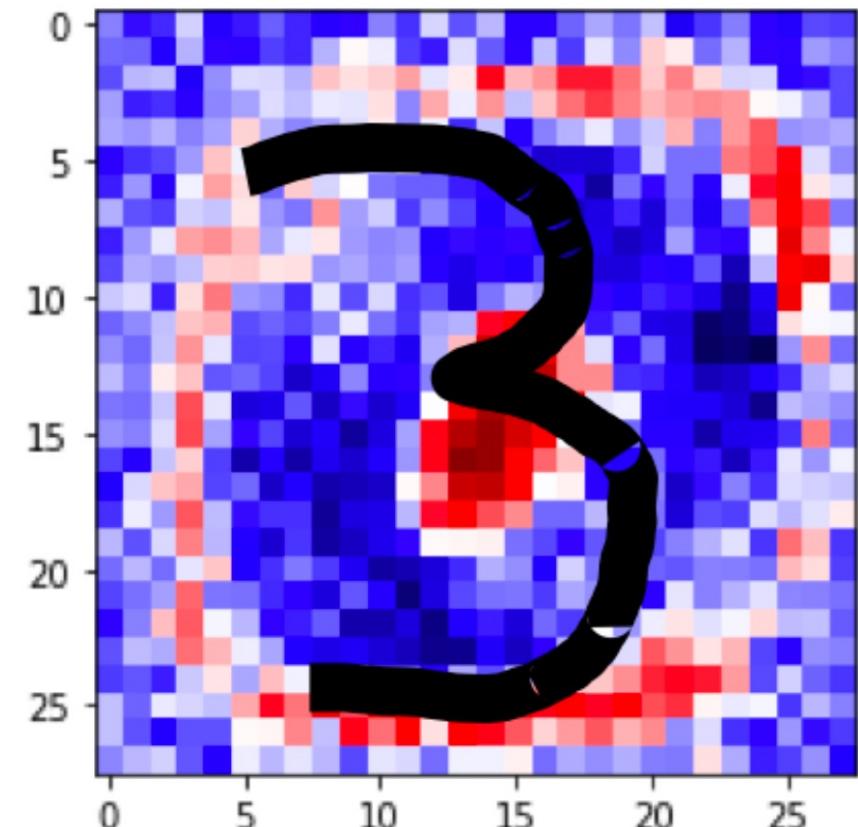
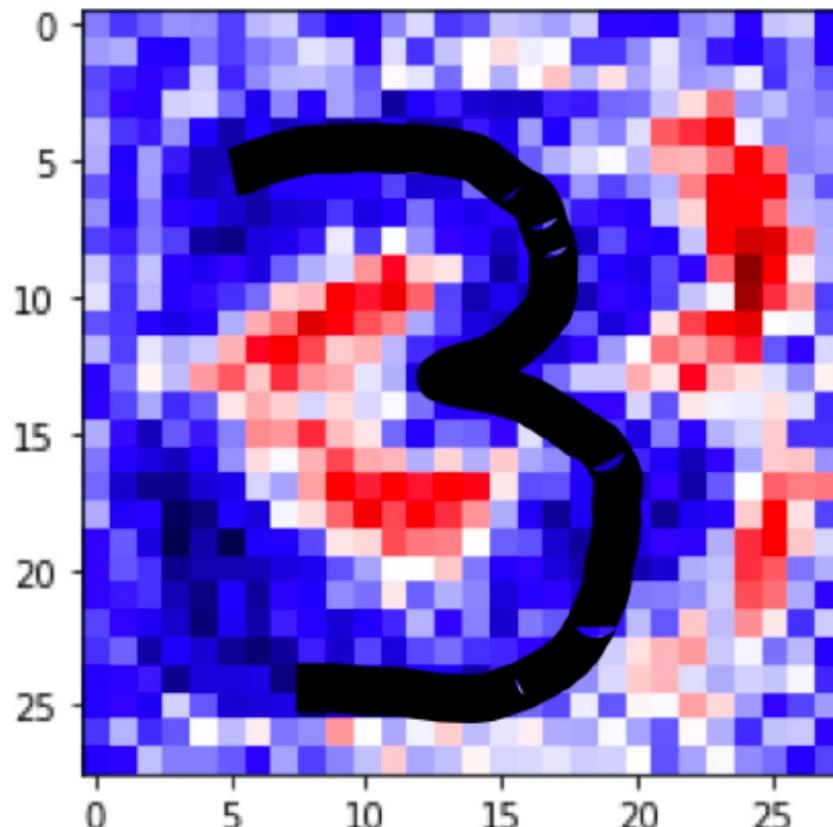
- To confirm that it is a good metric, let's imagine now that **we draw a three**
- it is clear that the red zone of the center of the previous model that we used for the zero will penalize the aforementioned metric since,
- as we can see in the figure, when writing a three we pass over



# Evidence of belonging

But, on the other hand, if the reference model is the one corresponding to number 3

we can see that the different possible traces that represent a three are mostly maintained in the blue zone.



# Probability of belonging

- The **second step** involves computing probabilities.
- Specifically **we turn the sum of evidences into predicted probabilities** using this function:

$$\text{Softmax}_i = \frac{e^{\text{evidence}_i}}{\sum_j e^{\text{evidence}_j}}$$

- softmax uses the exponential value of the calculated evidence and then normalizes them so that the sum equates to one, forming a probability distribution.

# Probability of belonging

$e^{evidence_i}$

- Intuitively, the effect obtained with the use of exponentials is that one more unit of evidence has a multiplier effect and one unit less has the inverse effect.
- The interesting thing about this function is that
  - **a good prediction** will have a single entry in the vector with a value close to 1, while the remaining entries will be close to 0.
  - in a **weak prediction**, there will be several possible labels, which will have more or less the same probability.

# Softmax

- A function that provides probabilities for each possible class in a **multi-class classification model**. The probabilities add up to exactly 1.0.
  - For example, softmax might determine that the probability of a particular image being a “7” at 0.8, a “9” at 0.17, ...

```
: predictions = model.predict(x_test)

: print(predictions[0])

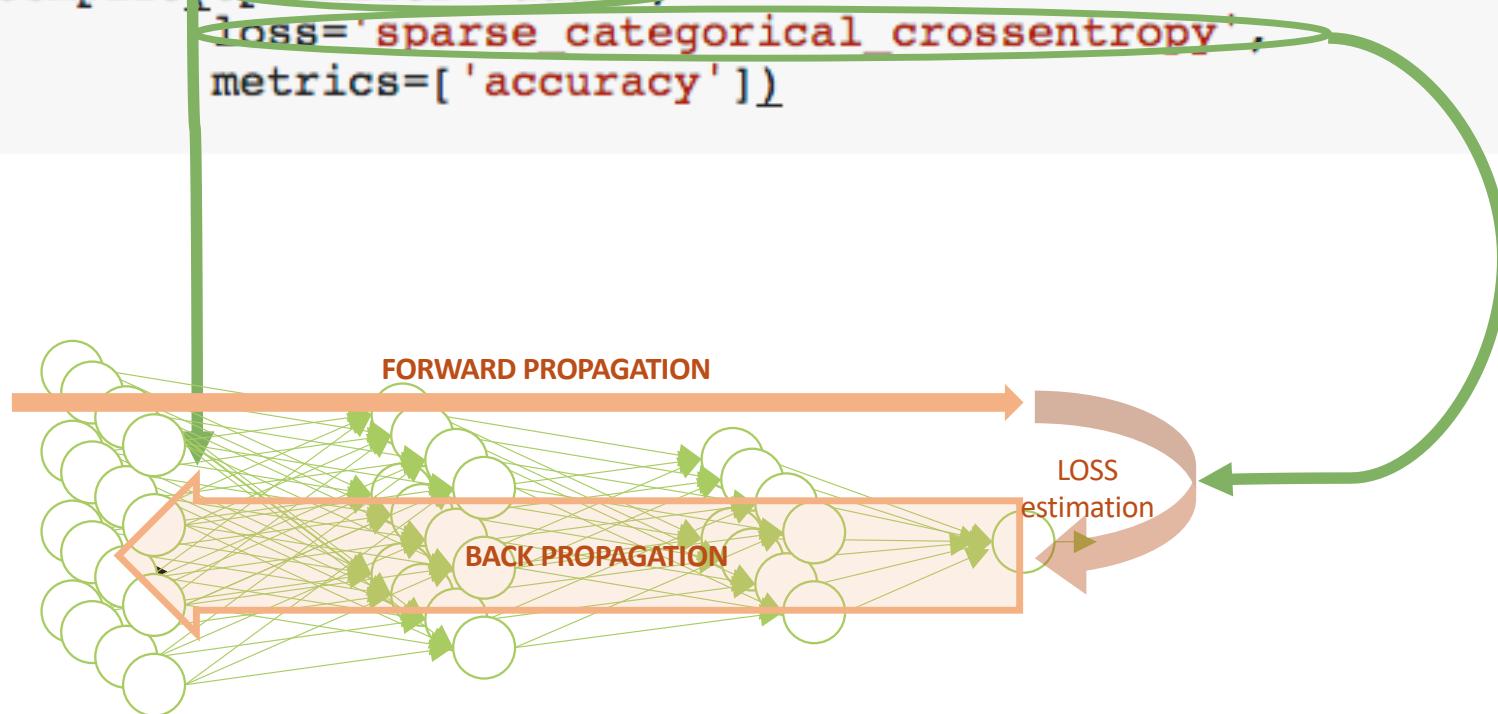
[0.    0.    0.    0.01 0.01 0.    0.    0.8   0.01 0.17]

: np.sum(predictions[0])

1.0
```

### 3. Compile model

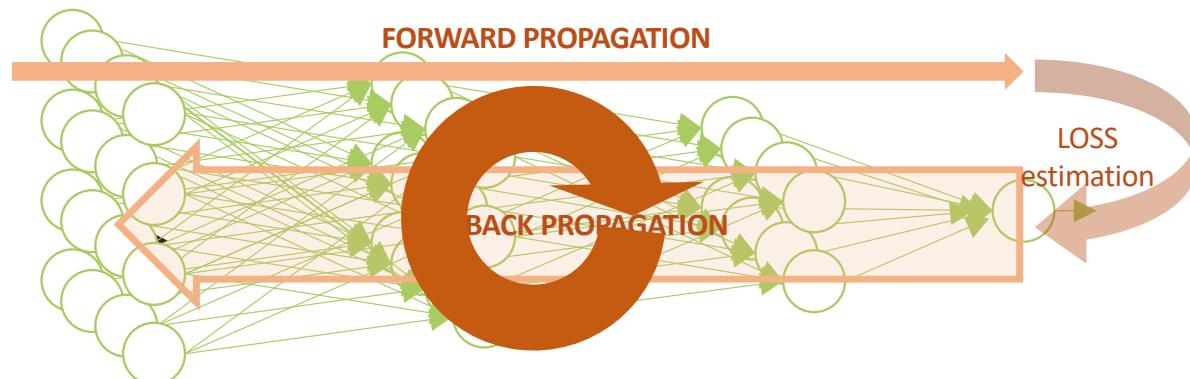
```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```



## 4. Fit model

```
[24] model.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5  
60000/60000 [=====] - 6s 98us/sample - loss: 0.4039 - accuracy: 0.8930  
Epoch 2/5  
60000/60000 [=====] - 6s 94us/sample - loss: 0.1958 - accuracy: 0.9441  
Epoch 3/5  
60000/60000 [=====] - 5s 88us/sample - loss: 0.1421 - accuracy: 0.9590  
Epoch 4/5  
60000/60000 [=====] - 5s 86us/sample - loss: 0.1101 - accuracy: 0.9688
```



## 5. Evaluate model

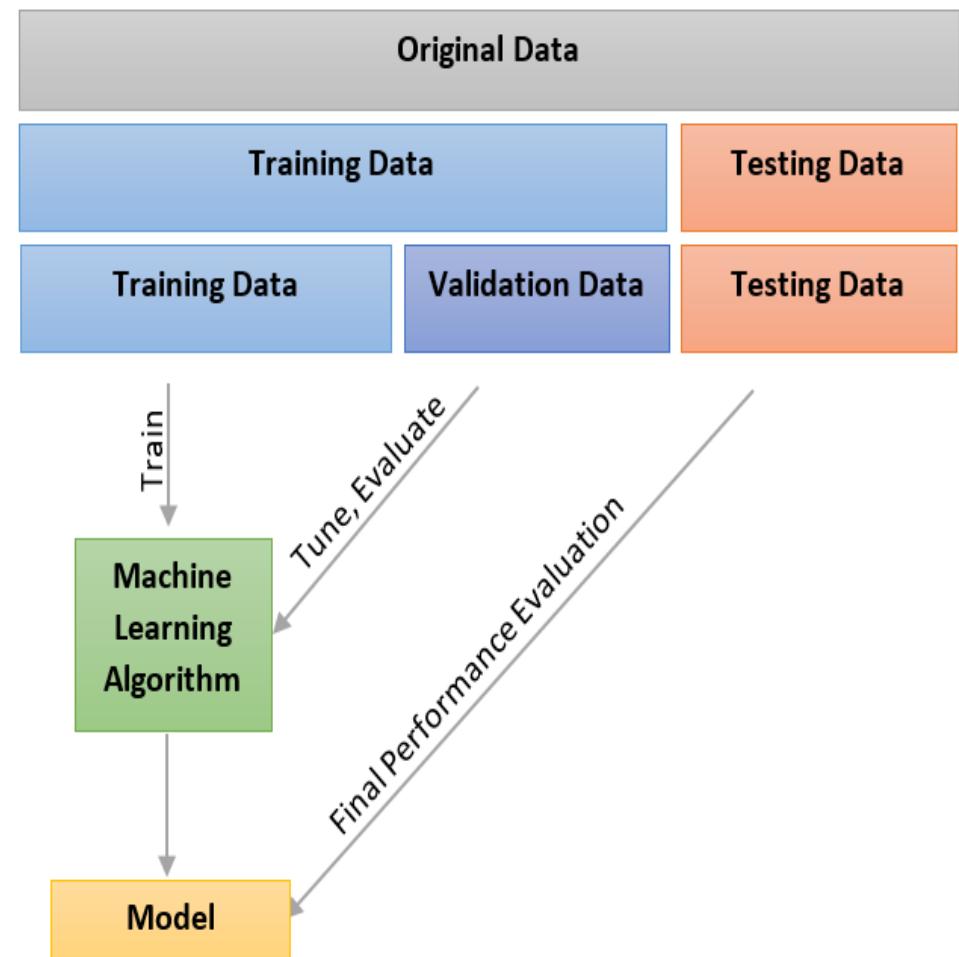
```
[25] test_loss, test_acc = model.evaluate(x_test, y_test)
    print ('Test accuracy:', test_acc)

→ 10000/10000 [=====] - 0s 48us/sample - loss: 0.0955 - accuracy: 0.9716
Test accuracy: 0.9716
```

## 5. Evaluate model

### Data

- Training set  
→ **For training**
- Validation set  
→ **For hyperparameter tuning**
- Test set  
→ **Test model performance**



## 5. Evaluate model

- Confusion matrix

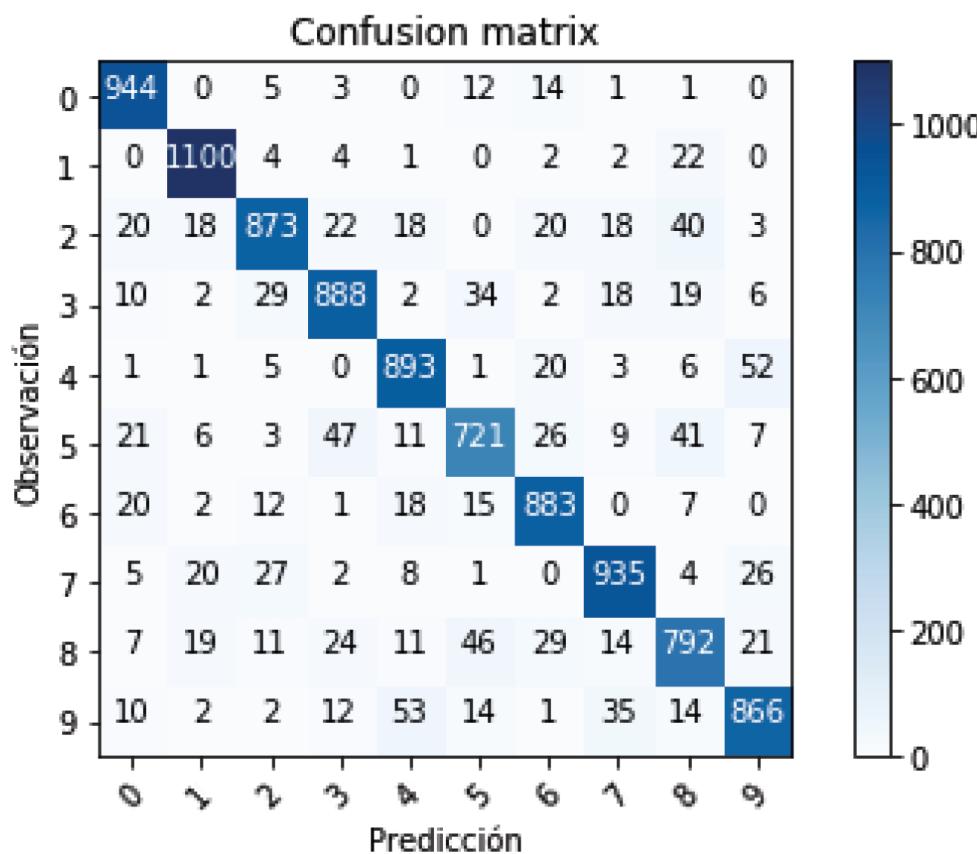
		Predicted class	
		positive	negative
Actual class	positive	TP	FN
	negative	FP	TN

- Accuracy

$$\text{Acc} = (\text{TP} + \text{TN}) / (\text{TP} + \text{FP} + \text{FN} + \text{TN})$$

## 5. Evaluate model

Confusion matrix for this model



## 6. Use the model

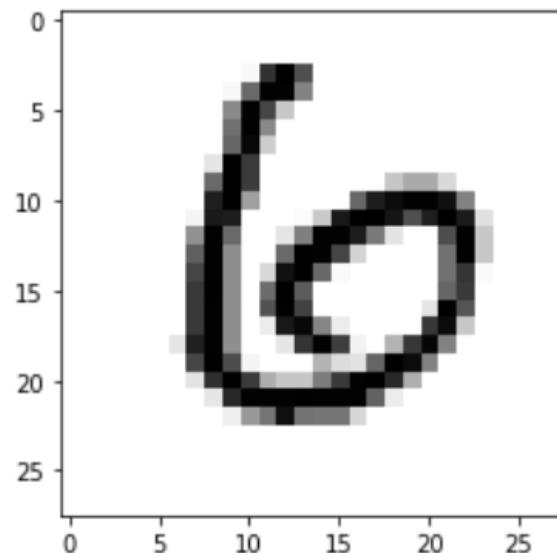
```
[31] predictions = model.predict(x_test)
     print (predictions)

[ [1.26404475e-05 2.49308573e-06 4.45205660e-04 ... 9.93939281e-01
  2.26089851e-05 3.87840439e-04]
  [1.40489414e-04 5.95517922e-03 9.92044926e-01 ... 7.36168673e-08
  8.24517483e-05 5.90428044e-08]
  [4.98437309e-08 9.97506320e-01 1.01395021e-03 ... 9.72060545e-04
  2.70858494e-04 1.64484100e-05]
  ...
  [1.58638017e-07 4.75961940e-07 7.93745471e-07 ... 1.92583975e-04
  1.82772157e-04 4.21503466e-03]
  [2.74540548e-06 5.57959866e-05 4.68355069e-07 ... 3.29460909e-06
  5.54329483e-03 1.57847660e-06]
  [2.11494735e-05 1.39356416e-06 7.07783474e-05 ... 2.89122028e-07
  5.49738729e-07 9.28026580e-08]]
```

## 6. Use the model

```
In [32]: x_test_old = x_test.reshape(10000, 28,28)  
plt.imshow(x_test_old[11], cmap=plt.cm.binary)
```

```
Out[32]: <matplotlib.image.AxesImage at 0x7ffec02fcb10>
```



## 6. Use the model

```
In [33]: predictions = model.predict(x_test)
```

```
In [34]: np.argmax(predictions[11])
```

Out[34]: 6

```
In [35]: print(predictions[11])
```

```
[0.0705287  0.01105603  0.19185546  0.00774726  0.05123473  0.03251035  
 0.536685    0.00300946  0.08143999  0.01393303]
```

```
In [36]: np.sum(predictions[11])
```

Out[36]: 1.0

# Pf: Deep Learning Basics (using Colab)

## ■ Tasks included:

- task 7.1 – Set up your google colab environment
- task 7.2 – Execute the provided notebook step-by-step
- task 7.3 – Improve the accuracy of your model

TODAY

part 1

- task 8.1 – Improving a basic cnn model
  - task 8.2 – Exporting the python script
  - task 8.3 – Running your first natural network on a login node
  - task 8.4 – Submitting your first gpu dl training with slurm
- 
- task 9.1 – Comparative implementation in pytorch and tensorflow

part 2

# Pf: Deep Learning Basics (using Colab)

## ■ In class:

- This lab is mainly **pedagogical**, aimed at helping you understand the fundamental concepts of deep learning.
- During lab class, the instructor will **assist each student individually** to consolidate understanding.

## ■ Deliverable & Evaluation:

- Upload a single PDF (per group) to the **racó@FIB** intranet.
- This PDF can simply include **screenshots or short notes** showing that you have completed each task.  
The goal is to confirm that you worked through the exercises, not to produce a polished “report”.

## ■ Evaluation day:

- A group will give a **brief and informal explanation** of their results.
- “No preparation is required” — **a few comments or screenshots are enough**.  
The focus is on learning, not on presentation style or timing.